

Paradigms in Temporal Localization: An Exhaustive Analysis of Classification-Based Video Moment Retrieval

The Landscape of Video Moment Retrieval and the Dominance of Regression

The proliferation of digital video content across online platforms, surveillance systems, and personal archives has created an urgent need for intelligent systems capable of fine-grained content analysis and retrieval.¹ While traditional video search systems focus on retrieving entire videos from a corpus, a more nuanced and increasingly critical task is Video Moment Retrieval (VMR), which aims to pinpoint specific temporal segments within a single, often lengthy and untrimmed, video that correspond to a user's query.² This task, also known by aliases such as Temporal Sentence Grounding in Videos (TSGV) and Natural Language Video Localization (NLVL), represents a significant challenge at the intersection of computer vision and natural language processing.³ The core objective is to, given a video

V and a natural language query Q , predict the precise start and end timestamps, (τ_s, τ_e) , of the moment described by Q .² This capability underpins a wide range of applications, from efficient video browsing and editing to advanced video question answering and summarization.²

The initial and most intuitive approaches to this problem were heavily influenced by established paradigms in computer vision, particularly object detection. These methods framed the localization of a temporal moment as a geometric regression problem, analogous to predicting the coordinates of a bounding box in an image. This section establishes this regression-based framework as the conventional paradigm, detailing its core mechanics, seminal models, and inherent limitations. Understanding this foundational approach is critical, as its shortcomings directly motivated the subsequent and profound shift towards treating temporal localization as a classification problem—a paradigm shift that is the central focus of

this report.

Defining the Task: From Video Search to Moment Retrieval

The VMR task is fundamentally more complex than related video understanding tasks, a distinction that clarifies its unique challenges. It differs from standard Video Retrieval (VR), which operates at a coarser granularity by retrieving entire videos from a large collection based on a query.³ VMR demands a much finer temporal precision, localizing a segment that might constitute only a small fraction of the source video's total duration.⁷

Furthermore, VMR is distinct from Temporal Action Localization (TAL). While TAL also identifies temporal boundaries of events, it is constrained to a predefined, closed set of action categories, such as "throwing" or "jumping".² For instance, a TAL system might be trained to find all instances of "LayUp Drill" from a vocabulary of known basketball moves.² In contrast, VMR leverages the descriptive power of natural language to handle an open set of complex, arbitrary activities that cannot be easily enumerated.² A query like "the small girl gets back on her feet after falling" describes a sequence of actions and their temporal relationship, a level of semantic complexity far beyond the scope of typical TAL systems.² This flexibility makes VMR a more general and powerful tool for video analysis but also introduces significant technical hurdles.

The primary challenges inherent to VMR can be summarized as follows:

1. **Cross-Modal Semantic Alignment:** The core difficulty lies in bridging the vast semantic gap between the continuous, high-dimensional visual data of a video and the discrete, symbolic nature of a language query.¹ The model must develop a shared understanding that links abstract linguistic concepts to concrete visual manifestations.
2. **Complex Query Comprehension:** Queries often describe intricate temporal relationships, ordering, and subtle interactions between objects and actions. For example, a query specifying the "second" occurrence of an event or an action that happens "after" another requires a model to perform temporal reasoning, not just simple content matching.²
3. **Temporal Search Space and Duration Variability:** Untrimmed videos can be extremely long, creating a vast temporal search space. The target moments themselves can have widely varying durations, from a few seconds to several minutes, and may be surrounded by long stretches of irrelevant content, making their localization a "needle in a haystack" problem.³
4. **Distinguishing Ambiguous Moments:** A video often contains multiple moments that are visually similar or share common objects and actions with the target moment. The model must discern the fine-grained semantic details specified in the query to distinguish the

correct moment from these visually plausible but incorrect distractors.⁹

Addressing these challenges has driven the evolution of VMR architectures, beginning with the adaptation of well-established computer vision techniques.

The Conventional Paradigm: Temporal Bounding Box Regression

The dominant early paradigm for VMR borrowed its conceptual framework directly from the field of object detection in 2D images. This intellectual lineage is not merely incidental; it shaped the very formulation of the problem. The task of finding a temporal segment on a 1D timeline was seen as a direct analogue to finding a 2D bounding box on an image plane. Consequently, the architectural solutions mirrored those of successful object detectors like Faster R-CNN, which operate on a proposal-then-regress pipeline. This approach typically involves two distinct stages: generating a set of candidate moments and then refining the boundaries of the best-matching candidate through regression.¹³

The first stage, **proposal generation**, involves creating a dense set of potential temporal segments that could contain the target moment. A common technique is the use of multi-scale temporal sliding windows, where windows of various predefined lengths are slid across the video's timeline to generate an exhaustive list of candidates.¹² Another approach uses predefined temporal anchors, which are reference segments of various scales and aspect ratios centered at different points in time, similar to the 2D anchors used in object detection.¹⁶

The second stage, **ranking and refinement**, evaluates these candidates. Each generated proposal is encoded into a feature vector, which is then compared against the feature vector of the language query to compute a cross-modal similarity or alignment score. The proposals are ranked based on this score. The highest-scoring proposal is selected as the most likely candidate, but its boundaries, derived from the coarse proposal generation step, are often imprecise. To address this, a regression head is employed. This head takes the features of the top-ranked proposal and the query as input and predicts a set of continuous-valued offsets. These offsets represent the adjustments needed to shift the candidate's start and end times to better align with the ground-truth moment. The final output is thus a refined temporal bounding box.

A seminal model that exemplifies this regression-based paradigm is the **Cross-modal Temporal Regression Localizer (CTRL)**.¹² The CTRL architecture clearly delineates the components of this approach. It uses a visual encoder to process candidate clips generated by sliding windows, a sentence encoder to process the query, and a multi-modal processing module to fuse these representations. This fused representation is then fed into a temporal

localization network, which is the heart of the regression mechanism. This network produces two sibling outputs: an alignment score, used for ranking the candidates, and a set of

clip location regression offsets. The model is trained to predict these continuous-valued offsets to refine the boundaries of the candidate clip. CTRL explicitly explored two types of regression targets: a parameterized offset that adjusts the clip's center point and length, and a non-parameterized offset that directly predicts the adjustments for the start and end points. The training objective for this component is typically a smooth L1 loss or a similar regression loss function, which penalizes the difference between the predicted and ground-truth offsets.¹²

While this proposal-and-regress paradigm provided a strong and intuitive starting point for the VMR task, its foundational reliance on object detection principles also meant it inherited their limitations. These drawbacks became increasingly apparent as researchers sought more efficient and accurate solutions, creating the necessary research pressure for a paradigm shift. The key limitations include:

- **Computational Inefficiency:** The generation of dense proposals via sliding windows or anchors is computationally expensive and memory-intensive. It results in an enormous number of candidate moments, the vast majority of which are irrelevant negative samples, leading to a highly imbalanced training problem and inefficient inference.¹⁶
- **Dependence on Hand-Crafted Components:** The performance of these models is highly sensitive to the choice of hand-designed components, such as the scales and aspect ratios of the sliding windows or anchors. These hyperparameters require careful, dataset-specific tuning and often fail to generalize to videos or moment types with characteristics not seen during training.²⁰
- **Error Propagation:** The two-stage nature of the pipeline is susceptible to error propagation. If the initial proposal generation stage fails to produce a candidate with sufficient overlap with the ground-truth moment, the subsequent regression head has no hope of recovering the correct boundaries, placing a hard upper bound on the model's performance.

The accumulation of these challenges—inefficiency, reliance on heuristics, and performance ceilings—directly motivated the research community to explore fundamentally different ways of formulating the localization problem. This led to the development of proposal-free, end-to-end architectures that could predict boundaries directly from the global video-query context. This search for a more elegant and effective solution paved the way for recasting temporal localization not as a geometric regression task, but as a classification problem, drawing inspiration from a different domain entirely: natural language processing.

Recasting Localization as Span Prediction: The

Question Answering Analogy

The limitations of the regression-based paradigm, inherited from object detection, spurred a significant conceptual shift in the VMR field. Researchers began to look beyond traditional computer vision for inspiration, finding a powerful and compelling analogy in the domain of natural language processing: span-based Question Answering (QA). This reframing—treating a video as a passage, a query as a question, and the target moment as the answer span—marked a pivotal moment of convergence between the fields and established the first major classification-based approach to temporal localization. This new perspective fundamentally changed how the problem was modeled, replacing geometric regression heads with probabilistic classifiers operating on sequences.

The Conceptual Leap: Video as Passage, Moment as Answer

The core idea behind this paradigm is a direct and elegant analogy: the task of finding a moment in a video that answers a language query is structurally identical to the task of finding a span of text in a document that answers a question.³ In span-based QA, a model is given a context passage (e.g., a Wikipedia article) and a question, and it must predict the start and end indices of the contiguous span of text within the passage that contains the answer. By mapping the components of VMR to this framework, the problem is transformed:

- The **untrimmed video** becomes the **context passage**.
- The **natural language query** becomes the **question**.
- The **target moment** becomes the **answer span**.

This conceptual leap had profound implications. It moved the problem from the domain of computer vision geometry, concerned with continuous coordinates and offsets, to the domain of NLP sequence processing, concerned with discrete tokens and indices. This allowed the VMR community to import and adapt the powerful and mature architectures that had been developed for extractive QA tasks, such as those benchmarked on the Stanford Question Answering Dataset (SQuAD). Instead of regressing four floating-point numbers for a bounding box, the task became classifying which of the discrete temporal units in a video sequence corresponded to the start and end of the answer. This shift from regression to classification offered a path toward more streamlined, end-to-end models that could bypass the cumbersome proposal generation stage entirely.

This cross-disciplinary transfer of methodology proved highly effective. It demonstrated that for certain complex, cross-modal reasoning tasks, paradigms developed for linguistic processing could be more suitable than those from traditional, vision-centric geometry. This

validation of an NLP-centric approach was not merely an incremental improvement; it represented a new way of thinking about temporal localization and prefigured the later, even deeper integration of language models into the core of vision-language architectures.

Architectural Deep Dive: VSLNet (Video Span Localizing Network)

The **Video Span Localizing Network (VSLNet)** stands as a prime exemplar of this QA-based approach, explicitly framing Natural Language Video Localization (NLVL) as a span prediction task.²² Its architecture is a direct adaptation of standard span-based QA models.

The VSLNet pipeline typically unfolds as follows:

1. **Embedding Layer:** The input video is first sampled into a sequence of discrete units (e.g., clips or frames), and a pre-trained visual feature extractor (like C3D or I3D) is used to generate a feature embedding for each unit. Similarly, the words in the language query are converted into textual embeddings. This step effectively discretizes the continuous video timeline into a sequence of tokens, making it amenable to sequence models.
2. **Contextual Encoder:** The sequence of video features and the sequence of query features are often processed by a bidirectional sequence model, such as a Bi-directional Long Short-Term Memory (Bi-LSTM) network or a Transformer encoder. This layer generates context-aware representations for each token, allowing the model to understand each video clip and query word in the context of its surroundings.
3. **Cross-Modal Interaction Layer:** A crucial component is a mechanism for modeling the fine-grained interaction between the video and the query. This is often achieved through a co-attention or bi-attention layer, which computes an affinity matrix between every video token and every query token. This allows the model to "attend" to the most relevant parts of the query when processing a specific part of the video, and vice-versa, producing a query-aware video representation.
4. **Prediction Head (The Classification Core):** This is where the paradigm shift from regression to classification is most evident. Instead of a regression head that outputs continuous values, the model employs a prediction head with two independent classifiers. After processing the query-aware video features, two separate linear layers, each followed by a softmax function, are applied to the sequence.
 - The first classifier outputs a probability distribution, P_{start} , over all temporal positions in the video sequence. $P_{start}(i)$ represents the model's confidence that the i -th temporal unit is the start of the target moment.
 - The second classifier outputs another probability distribution, P_{end} , over all temporal positions. $P_{end}(j)$ represents the confidence that the j -th temporal unit is the end of the target moment.

The final moment localization is determined not by regressing offsets but by selecting the

indices. The model searches for the pair of indices (i, j) where $i \leq j$ that maximizes the joint probability score, typically calculated as $P_{\text{start}}(i) \times P_{\text{end}}(j)$. This formulation transforms the localization task into two distinct T-way classification problems, where T is the number of temporal units in the video sequence.

The training objective reinforces this classification-based nature. The model is optimized using a cross-entropy loss function, which is the standard loss for multi-class classification. The total loss is the sum of the negative log-likelihoods for the ground-truth start index and the ground-truth end index.²³ This directly encourages the model to assign high probabilities to the correct start and end positions and low probabilities to all others.

While the QA analogy is powerful, VSLNet's authors also recognized the inherent differences between the modalities. Video is a continuous, temporally coherent medium, whereas language is discrete and structured syntactically.²³ To bridge this gap, VSLNet introduced a

Query-Guided Highlighting (QGH) strategy. This mechanism first predicts a coarse, "highlighted" region of the video that is likely to contain the answer span, effectively reducing the search space and allowing the main prediction head to focus its attention on the most relevant temporal context.²³

This shift to span prediction, however, introduces a critical and often implicit trade-off. By discretizing the continuous video timeline into a sequence of features, the model's temporal resolution becomes fundamentally tied to the sampling rate. A higher sampling rate yields a more precise localization potential but results in a longer sequence. For architectures like the Transformer, where the computational complexity of self-attention is quadratic with respect to sequence length ($O(T^2)$), this creates a severe bottleneck. A lower sampling rate is more computationally feasible but risks "blurring" the specific boundary of the target moment, as query-relevant frames may be skipped or averaged out with irrelevant ones.⁷ This trade-off between temporal precision and computational cost is a fundamental constraint of classification-based span prediction methods and a key consideration in their design and application.

Frame-Level Classification and Alternative Formulations

While the span-prediction framework offered a powerful, end-to-end alternative to regression, it is not the only way to formulate temporal localization as a classification problem. Other approaches have emerged that deconstruct the task in different ways, focusing on making independent classifications at the level of individual frames or proposing conceptual

frameworks that decouple semantic understanding from temporal grounding. These methods share the common thread of replacing a single, complex regression or span-prediction task with a series of simpler, more localized classification decisions. This move towards deconstruction represents a significant engineering principle: breaking down a multifaceted problem into more manageable sub-problems can lead to more robust, interpretable, and flexible models.

Probabilistic Boundary Prediction

An alternative to predicting a single, contiguous span is to adopt a "bottom-up" approach that makes probabilistic predictions for each temporal unit (frame or clip) independently.⁷ In this formulation, the model does not directly output a start-end pair. Instead, it processes the fused video-query features and, for each timestamp

t in the video, performs one or more binary or multi-class classification tasks.

The typical outputs of such a model are frame-wise probability scores, such as:

- **Start Probability $P(\text{is_start} \mid \text{frame}_t)$:** The probability that frame t is the starting boundary of the target moment.
- **End Probability $P(\text{is_end} \mid \text{frame}_t)$:** The probability that frame t is the ending boundary of the target moment.
- **Saliency/Inclusion Probability $P(\text{is_inside} \mid \text{frame}_t)$:** The probability that frame t falls within the temporal span of the target moment. This is often used to identify the general region of interest.

After the network produces these probability sequences for the entire video, the final moment boundaries are not directly predicted but are instead inferred through a post-processing step. This step typically involves applying a threshold to the probability maps to identify candidate start and end points and then using a grouping or matching algorithm to pair them into valid temporal segments.¹⁸ For example, for each candidate start point with a high probability, the algorithm might search for a subsequent end point that also has a high probability, with the final moment selected based on a combined confidence score.

The Hierarchical Uncertainty-based Active Learning (HUAL) framework, while primarily designed to reduce annotation costs, is built upon this very principle of frame-level classification.¹⁹ In its interactive loop, a weakly-supervised model first generates initial binary pseudo-labels for each frame (classifying it as relevant or not). A fully-supervised model is then trained on start and end boundary probability distributions that are derived from these sparse, binary frame-level annotations. This demonstrates the fundamental utility of treating each frame as a sample for classification, where the collective classifications are then used to

construct the final temporal localization. This approach provides dense supervision signals and can be more resilient to the slight inaccuracies or "soft boundaries" common in human annotations, as it does not force a single, hard decision for the entire span at once.

The "Classification-then-Grounding" Framework: A Conceptual Paradigm

A more radical deconstruction of the VMR task is offered by the "classification-then-grounding" framework, a conceptual paradigm introduced in the context of Video Scene Graph Generation (VidSGG).²⁶ Although not originally designed for VMR, its principles are directly applicable and offer a powerful alternative to monolithic localization models. The framework explicitly separates the problem into two distinct, sequential stages:

1. **Classification Stage:** The first stage is purely concerned with semantic understanding. It ignores the temporal localization aspect and focuses on classifying *what* is happening. In the context of VidSGG, this involves identifying all object entities (e.g., "person," "motorcycle") and the predicate relationships between them (e.g., "ride") that occur anywhere in the video.
2. **Grounding Stage:** The second stage takes the semantic concepts identified by the classification stage and performs temporal grounding. For each classified concept (e.g., the semantic triplet <person, ride, motorcycle>), it treats this concept as a new, more structured query and runs a dedicated temporal localization model to find *when* this specific event occurs in the video.

This framework could be conceptually applied to VMR in a straightforward manner.²⁶ The

classification stage would involve parsing the natural language query to identify its core semantic components—the subjects, objects, actions, and their relationships. For the query "a person riding a motorcycle," this stage would classify "person," "riding," and "motorcycle" as the key semantic elements. The **grounding stage** would then take these classified elements as input and search the video for temporal segments where a "person" is visually present and engaged in the action of "riding" a "motorcycle."

The primary advantage of this decoupled approach is that it disentangles the complex, intertwined challenges of multimodal semantic reasoning and precise temporal localization. By first establishing a robust understanding of the query's meaning (the "what"), the subsequent grounding task (the "when") becomes significantly more constrained and targeted. This modularity avoids the difficulty of training a single, end-to-end model to solve both problems simultaneously.

The increasing prevalence and power of large-scale, pre-trained foundation models make this

"classification-then-grounding" paradigm more viable and attractive than ever before. The classification stage, which requires a deep semantic understanding of both language and visual content, is a task for which modern Multimodal Large Language Models (MLLMs) are exceptionally well-suited.¹ These models, pre-trained on vast web-scale datasets, possess powerful zero-shot classification and reasoning capabilities that can be leveraged to parse the query and identify relevant visual concepts without requiring extensive task-specific training.²⁸ A large MLLM could therefore serve as a powerful "semantic front-end" for the classification stage. The grounding stage could then be handled by a more specialized, lightweight temporal localization model, which receives the structured semantic output from the MLLM as a much cleaner, less ambiguous input than the raw natural language query. This modular design, leveraging the distinct strengths of different classes of models, represents a promising direction for building more accurate, interpretable, and efficient VMR systems.

The Generative Turn: Regression-Free Localization with MLLMs

The evolution of temporal localization paradigms has recently entered a new, transformative phase driven by the remarkable capabilities of Multimodal Large Language Models (MLLMs). This latest shift recasts VMR not as a regression or discriminative classification task, but as an open-ended, generative sequence-to-sequence problem.²⁷ In this framework, the model does not predict numerical coordinates or classify temporal indices; instead, it generates a natural language string that explicitly contains the predicted timestamps. This approach is inherently regression-free and represents the ultimate culmination of treating localization as a classification problem—where each generated token, including the digits and symbols that form the timestamps, is the result of a classification over the model's vast vocabulary. This generative turn effectively subsumes the entire VMR task within the language modeling paradigm, marking a profound change in how vision-language problems are conceptualized and solved.

VMR as a Sequence-to-Sequence Task

The new paradigm fundamentally alters the input-output format of the VMR task. Instead of designing specialized prediction heads (for regression or classification), the model leverages a general-purpose, pre-trained generative MLLM. The task is framed as a prompted question-answering session:

- **Input:** The model is provided with a carefully constructed input sequence that combines the natural language query with visual information from the video. This typically involves sampling a sparse set of frames from the video and interleaving their visual features with the text query in a single, unified sequence.²⁷
- **Output:** The model's task is to auto-regressively generate a textual response that answers the query by providing the relevant timestamps. For example, given a query, the model might generate the string: "The moment you are looking for occurs between 00:15.3 and 00:28.1".²⁷

This formulation is, at its core, a series of classification decisions. At each step of the generation process, the language model predicts a probability distribution over its entire vocabulary (which includes words, numbers, punctuation, etc.) and selects the next token. The temporal boundaries emerge not from a specialized geometric or indexing mechanism but from the model's ability to generate the correct sequence of character tokens that represent the timestamps. This approach completely obviates the need for any task-specific architectural components like regression heads or boundary classifiers, leading to a much simpler and more unified model design.

This shift from discriminative to generative modeling is significant. It aligns VMR with the broader trend in artificial intelligence towards leveraging large foundation models for in-context learning and instruction-following.¹ The model is not a specialized VMR solver but a general-purpose vision-language reasoner that performs VMR when prompted to do so. This opens the door to greater flexibility, as the same base model could potentially perform a variety of video understanding tasks—such as VMR, dense video captioning, or video question answering—simply by altering the input prompt, without any change to the model architecture itself.

Architectural Deep Dive: Mr. BLIP

A state-of-the-art model that epitomizes this generative paradigm is **Mr. BLIP**.²⁷ Its design and methodology highlight the key principles of this new approach.

Model Overview: Mr. BLIP is a multimodal, single-stage model that demonstrates the surprising effectiveness of using an image-text pre-trained MLLM for the video task of moment retrieval. Its defining characteristics include:

- **Foundation Model-Based:** It is built directly upon a pre-trained MLLM, specifically BLIP-2. This is a critical design choice, as it allows the model to inherit the powerful visual comprehension and language reasoning capabilities learned from massive web-scale image-text corpora, without requiring any expensive, large-scale video-language

pre-training.²⁷

- **Simple and Unified Architecture:** The model avoids the complex, highly specialized architectures of many prior VMR methods. It does not require additional input signals like audio or ASR transcripts and consists of a single, end-to-end generative model without separate heads for different sub-tasks.²⁷

Input Representation and Implicit Temporal Reasoning: The most innovative aspect of Mr. BLIP lies in its input representation. To enable the MLLM to reason about time, the model is fed a novel multimodal sequence where visual frames and temporal information are explicitly interleaved. The process is as follows:

1. A fixed number of frames are sparsely sampled from the input video.
2. A prompt is constructed that includes the natural language query.
3. The sampled frames are then inserted into this prompt, and crucially, each frame is accompanied by a special text token representing its absolute timestamp (e.g., <image> A frame at 10.5 seconds. <image> A frame at 12.0 seconds....).

This interleaved sequence of text, image features, and timestamp tokens is then fed as a single input to the MLLM. The model's self-attention mechanism is thus able to directly process and learn the relationships between the visual content of a frame, its explicit temporal position in the video, and the semantics of the language query.²⁷ The temporal localization is not performed by a dedicated module; rather, it emerges implicitly from the MLLM's ability to reason over this rich, contextually grounded sequence and generate a textual answer containing the correct timestamps.

Advantages of the Generative Approach: This method offers several significant advantages over previous paradigms:

- **Flexibility in Prediction:** Unlike proposal-based methods that are constrained to a fixed number of predictions, the generative model can output a variable number of temporal intervals in its response. This is particularly useful for queries that may correspond to multiple distinct moments in a video, as the model can simply list all of them in its generated text.²⁷
- **Leveraging Pre-trained Knowledge:** By building on a powerful, pre-trained MLLM, the model can perform complex reasoning that goes beyond simple pattern matching. It can leverage the "world knowledge" embedded in the MLLM's parameters to better understand the query and its visual correspondence.
- **Architectural Simplicity:** The entire localization process is unified within the generative language model, eliminating the need for separate regression or classification heads and complex post-processing steps like non-maximum suppression.

The success of this generative approach completes the paradigm shift initiated by the QA analogy. It does not merely borrow a tool from NLP; it fully subsumes the VMR problem within the language modeling framework. In this view, visual data is simply another form of

information—another type of token—to be processed in a sequence. The future of VMR and other vision-language tasks may therefore lie less in designing novel, task-specific architectures and more in discovering increasingly effective ways to represent visual and temporal information as a "language" that large foundation models can natively comprehend.

Synthesis, Comparative Analysis, and Future Trajectories

The evolution of temporal localization in Video Moment Retrieval (VMR) reveals a clear and consistent trajectory away from traditional regression-based methods towards more sophisticated, classification-centric paradigms. This progression—from geometric regression to span-based classification, and ultimately to token-by-token generative modeling—is not merely a series of architectural updates but reflects a fundamental reconceptualization of the task itself. Each paradigm offers a distinct set of strengths and weaknesses, and their comparative analysis illuminates the current state of the field and points toward future research directions. The overarching trend is one of increasing abstraction, flexibility, and deeper integration with the powerful reasoning capabilities of large-scale language models.

Comparative Analysis of Localization Paradigms

The different approaches to temporal localization in VMR can be categorized based on their core mechanism for predicting start and end boundaries. Each paradigm represents a different set of assumptions about the nature of the problem and involves distinct architectural trade-offs.

Paradigm	Core Mechanism	Key Exemplar Models	Strengths	Limitations	Key Citations
Temporal Bounding Box Regression	Predicts continuous -valued start/end times or	CTRL	- Conceptually aligned with traditional	- Computationally inefficient due to	¹²

	offsets relative to candidate proposals (anchors/windows).		CV object detection. - Can achieve high precision with good proposals.	dense proposals. - Relies on hand-designed components (anchors, NMS). - Prone to error propagation in two-stage pipelines.	
Span Prediction as Classification	Predicts two probability distributions over a discrete sequence of temporal units for start and end indices.	VSLNet	- End-to-end and proposal-free. - Elegant formulation leveraging mature NLP (QA) architectures. - Computationally more efficient than dense proposal methods.	- Performance is sensitive to the temporal discretization rate. - Fixed-length sequence inputs can struggle with very long videos.	22
Frame-Level Classification	Predicts the probability of each frame being a start, end, or internal	(Component in) SeqPAN, HUAL	- Provides dense, frame-level supervision. - Can be more robust to	- Requires post-processing to assemble moments from probability	18

	point of the moment.		noisy boundary annotations . - Conceptually simple and modular.	maps. - Can produce fragmented or inconsistent predictions without proper regularization.	
Generative Sequence-to-Sequence	Generates a textual string representing the start and end timestamps as a sequence of tokens.	Mr. BLIP, UniTime	- Leverages immense power of pre-trained MLLMs. - Highly flexible (e.g., variable number of outputs). - Enables zero-shot/few-shot performance and complex reasoning. - Architecturally simple (no task-specific heads).	- High computational cost for inference. - Limited by MLLM context windows for long videos. - Potential for factual inaccuracies (hallucination).	27

The journey across these paradigms reveals a fascinating evolution in how a "temporal coordinate" is represented. In the **regression** framework, a coordinate is a continuous floating-point number in a 1D geometric space. In the **span-prediction classification** framework, a coordinate is a discrete index within a finite sequence. Finally, in the **generative MLLM** framework, a coordinate is a symbolic string of text, constructed from a vocabulary of

tokens. This progression from a mathematical to a symbolic and ultimately linguistic representation of time suggests that to achieve human-like understanding of video based on language, the model's internal reasoning and output format should themselves be more language-like. This conceptual arc provides the strongest justification for why classification-based and generative approaches have become dominant and represent the future of the field.

Emerging Trends and Future Trajectories

The analysis of these paradigms illuminates several key trends that are shaping the future of VMR and related temporal grounding tasks.

1. **The Dominance of Foundation Models:** There is an undeniable and accelerating shift away from training specialized models from scratch on limited VMR datasets. The field is now heavily centered on adapting large, pre-trained vision-language foundation models, such as CLIP and especially MLLMs.¹ These models provide a powerful starting point, imbuing VMR systems with rich semantic priors and reasoning capabilities learned from web-scale data, which significantly improves performance and generalization.
2. **The Pursuit of Unified and Multi-Task Frameworks:** Modern architectures, particularly those based on the Transformer, are increasingly designed as unified frameworks capable of jointly solving multiple related tasks. For example, models like Moment-DETR and UVCOM are designed to perform both Video Moment Retrieval and Highlight Detection within a single, shared architecture.²² This trend reflects a move towards more general-purpose video understanding systems rather than narrow, single-task solvers.
3. **Reducing Annotation Dependency:** The high cost and subjectivity of annotating precise temporal boundaries have fueled a strong interest in learning paradigms that reduce this dependency. This includes research into weakly-supervised methods that learn from video-sentence pairs without timestamps¹¹, zero-shot methods that leverage foundation models to perform VMR without any task-specific fine-tuning³⁵, and self-supervised pre-training on large corpora of unlabeled videos to learn robust representations before fine-tuning on downstream tasks.³⁷ Classification and generative frameworks are particularly well-suited to these low-supervision settings.

Based on these trends, several critical research questions emerge for the future:

- **Scaling to Long-Form Video:** A primary bottleneck for MLLM-based generative models is their limited context window, which makes processing hour-long videos computationally infeasible. Future research will need to focus on developing more efficient video representations and processing strategies. Promising directions include hierarchical or recursive approaches, where a model first identifies coarse segments of interest and then "zooms in" for fine-grained localization, as explored by models like

ReVisionLLM.¹

- **Explicit "Classification-then-Grounding" Architectures:** While the conceptual "classification-then-grounding" framework is compelling, few works have explicitly implemented it using modern foundation models. A promising research avenue would be to design modular systems where a powerful MLLM acts as a "semantic parser" to deconstruct a query, passing a structured representation to a separate, efficient temporal grounding module. This could lead to more interpretable, controllable, and potentially more accurate VMR systems.
- **Addressing Data Bias and Ambiguity:** VMR datasets are known to contain inherent biases and ambiguities, such as spurious correlations between queries and video backgrounds, which can lead models to learn superficial shortcuts rather than true semantic understanding.⁴ Future work must investigate how to mitigate these biases within the new classification and generative paradigms, perhaps through advanced data augmentation techniques, causal intervention methods, or by designing prompts for MLLMs that encourage more robust reasoning.

In conclusion, the field of Video Moment Retrieval has undergone a profound transformation, moving decisively from regression-based localization to a spectrum of classification-based approaches. This shift, driven by the search for greater efficiency, accuracy, and flexibility, has culminated in the adoption of generative MLLMs that treat VMR as a language modeling task. This final paradigm, while computationally demanding, offers unprecedented capabilities in complex reasoning and flexible prediction, firmly establishing that the path to deeper video understanding is paved with the principles of classification and language.

Works cited

1. (PDF) A Survey on Video Temporal Grounding with Multimodal Large Language Model, accessed August 28, 2025, https://www.researchgate.net/publication/394524600_A_Survey_on_Video_Temporal_Grounding_with_Multimodal_Large_Language_Model
2. Video Moment Retrieval to Capture the Highlighting Moments in Video (Part.1) - Medium, accessed August 28, 2025, <https://medium.com/@myschang/video-moment-retrieval-to-capture-the-highlighting-moments-in-video-part-1-4bd7ede78c93>
3. Video Moment Retrieval, accessed August 28, 2025, <https://www3.ntu.edu.sg/home/axsun/slides/VideoMomentRetrieval.pdf>
4. Temporal Sentence Grounding in Videos: A Survey and Future Directions - ResearchGate, accessed August 28, 2025, https://www.researchgate.net/publication/369339968_Temporal_Sentence_Grounding_in_Videos_A_Survey_and_Future_Directions
5. Soldelli/Awesome-Temporal-Language-Grounding-in-Videos - GitHub, accessed August 28, 2025, <https://github.com/Soldelli/Awesome-Temporal-Language-Grounding-in-Videos>
6. Video Corpus Moment Retrieval with Contrastive Learning - arXiv, accessed

- August 28, 2025, <https://arxiv.org/pdf/2105.06247>
7. Fewer Steps, Better Performance: Efficient Cross-Modal Clip Trimming for Video Moment Retrieval Using Language - AAAI Publications, accessed August 28, 2025, <https://ojs.aaai.org/index.php/AAAI/article/view/27941/27902>
 8. A Parallel Transformer Framework for Video Moment Retrieval - DORAS | DCU Research Repository, accessed August 28, 2025, <https://doras.dcu.ie/30558/2/3652583.3658096.pdf>
 9. Background-Aware Moment Detection for Video Moment Retrieval - CVF Open Access, accessed August 28, 2025, https://openaccess.thecvf.com/content/WACV2025/papers/Jung_Background-Aware_Moment_Detection_for_Video_Moment_Retrieval_WACV_2025_paper.pdf
 10. Open-Vocabulary Temporal Action Localization using Multimodal Guidance - BMVA Archive, accessed August 28, 2025, https://bmva-archive.org.uk/bmvc/2024/papers/Paper_1013/paper.pdf
 11. Summarization of weakly-supervised and unsupervised video moment localization approaches. - ResearchGate, accessed August 28, 2025, https://www.researchgate.net/figure/Summarization-of-weakly-supervised-and-unsupervised-video-moment-localization-approaches_tbl2_371537174
 12. TALL: Temporal Activity Localization via ... - CVF Open Access, accessed August 28, 2025, https://openaccess.thecvf.com/content_ICCV_2017/papers/Gao_TALL_Temporal_Activity_ICCV_2017_paper.pdf
 13. A Survey on Video Moment Localization - arXiv, accessed August 28, 2025, <https://arxiv.org/pdf/2306.07515>
 14. Survey on Video Moment Retrieval - SciEngine, accessed August 28, 2025, <https://www.sciengine.com/doi/10.13328/j.cnki.jos.006707>
 15. Multi-Modal Relational Graph for Cross-Modal Video Moment Retrieval - CVF Open Access, accessed August 28, 2025, https://openaccess.thecvf.com/content/CVPR2021/papers/Zeng_Multi-Modal_Relational_Graph_for_Cross-Modal_Video_Moment_Retrieval_CVPR_2021_paper.pdf
 16. MomentDiff: Generative Video Moment Retrieval from Random to Real - arXiv, accessed August 28, 2025, <https://arxiv.org/pdf/2307.02869>
 17. Fast Video Moment Retrieval - CVF Open Access, accessed August 28, 2025, https://openaccess.thecvf.com/content/ICCV2021/papers/Gao_Fast_Video_Moment_Retrieval_ICCV_2021_paper.pdf
 18. Hierarchical Local-Global Transformer for Temporal Sentence Grounding - arXiv, accessed August 28, 2025, <https://arxiv.org/pdf/2208.14882>
 19. Are Binary Annotations Sufficient? Video ... - CVF Open Access, accessed August 28, 2025, https://openaccess.thecvf.com/content/CVPR2023/papers/Ji_Are_Binary_Annotations_Sufficient_Video_Moment_Retrieval_via_Hierarchical_Uncertainty-Based_CVPR_2023_paper.pdf
 20. Learning Salient Boundary Feature for Anchor-free Temporal Action ..., accessed August 28, 2025, <https://arxiv.org/abs/2103.13137>
 21. End-to-End Object Detection with Transformers, accessed August 28, 2025,

- https://www.ecva.net/papers/eccv_2020/papers_ECCV/papers/123460205.pdf
22. MS-DETR: Towards Effective Video Moment Retrieval and Highlight Detection by Joint Motion-Semantic Learning - ResearchGate, accessed August 28, 2025, https://www.researchgate.net/publication/393771456_MS-DETR_Towards_Effective_Video_Moment_Retrieval_and_Highlight_Detection_by_Joint_Motion-Semantic_Learning
 23. Span-based Localizing Network for Natural ... - ACL Anthology, accessed August 28, 2025, <https://aclanthology.org/2020.acl-main.585.pdf>
 24. [2102.13558] Natural Language Video Localization: A Revisit in Span-based Question Answering Framework - arXiv, accessed August 28, 2025, <https://arxiv.org/abs/2102.13558>
 25. Span-based Localizing Network for Natural Language Video Localization - ACL Anthology, accessed August 28, 2025, <https://aclanthology.org/2020.acl-main.585/>
 26. Classification-Then-Grounding: Reformulating Video Scene Graphs As Temporal Bipartite Graphs - CVF Open Access, accessed August 28, 2025, https://openaccess.thecvf.com/content/CVPR2022/papers/Gao_Classification-Then-Grounding_Reformulating_Video_Scene_Graphs_As_Temporal_Bipartite_Graphs_CVPR_2022_paper.pdf
 27. The Surprising Effectiveness of Multimodal Large Language Models for Video Moment Retrieval - arXiv, accessed August 28, 2025, <https://arxiv.org/html/2406.18113v2>
 28. The Surprising Effectiveness of Multimodal Large Language Models for Video Moment Retrieval - ResearchGate, accessed August 28, 2025, https://www.researchgate.net/publication/381736975_The_Surprising_Effectiveness_of_Multimodal_Large_Language_Models_for_Video_Moment_Retrieval
 29. A Survey on Video Temporal Grounding with Multimodal Large Language Model - arXiv, accessed August 28, 2025, <https://arxiv.org/abs/2508.10922>
 30. [2506.18883] Universal Video Temporal Grounding with Generative Multi-modal Large Language Models - arXiv, accessed August 28, 2025, <https://arxiv.org/abs/2506.18883>
 31. VideoGrounding-DINO: Towards Open-Vocabulary Spatio-Temporal Video Grounding, accessed August 28, 2025, https://openaccess.thecvf.com/content/CVPR2024/papers/Wasim_VideoGrounding-DINO_Towards_Open-Vocabulary_Spatio-Temporal_Video_Grounding_CVPR_2024_paper.pdf
 32. Bridging the Gap: A Unified Video Comprehension Framework for Moment Retrieval and Highlight Detection - CVF Open Access, accessed August 28, 2025, https://openaccess.thecvf.com/content/CVPR2024/papers/Xiao_Bridging_the_Gap_A_Unified_Video_Comprehension_Framework_for_Moment_CVPR_2024_paper.pdf
 33. UMT: Unified Multi-Modal Transformers for Joint Video Moment Retrieval and Highlight Detection - CVF Open Access, accessed August 28, 2025, https://openaccess.thecvf.com/content/CVPR2022/papers/Liu_UMT_Unified_Multi-Modal_Transformers_for_Joint_Video_Moment_Retrieval_and_CVPR_2022_pape

[r.pdf](#)

34. Weakly-Supervised Video Moment Retrieval via Semantic Completion Network - AAAI, accessed August 28, 2025, <https://cdn.aaai.org/ojs/6820/6820-13-10049-1-10-20200524.pdf>
35. Zero-shot Video Moment Retrieval With Off-the-Shelf Models - Proceedings of Machine Learning Research, accessed August 28, 2025, <https://proceedings.mlr.press/v203/diwan23a/diwan23a.pdf>
36. A Lightweight Moment Retrieval System with Global Re-Ranking and Robust Adaptive Bidirectional Temporal Search - arXiv, accessed August 28, 2025, <https://arxiv.org/html/2504.09298v1>
37. Vid-Morp: Video Moment Retrieval Pretraining from Unlabeled Videos in the Wild - arXiv, accessed August 28, 2025, <https://arxiv.org/abs/2412.00811>
38. ReVisionLLM: Recursive Vision-Language Model for Temporal Grounding in Hour-Long Videos, accessed August 28, 2025, https://openaccess.thecvf.com/content/CVPR2025/papers/Hannan_ReVisionLLM_Recursive_Vision-Language_Model_for_Temporal_Grounding_in_Hour-Long_Videos_CVPR_2025_paper.pdf
39. Reverse Distribution Based Video Moment Retrieval for Effective Bias Elimination, accessed August 28, 2025, <https://ojs.aaai.org/index.php/AAAI/article/download/33302/35457>
40. The Devil is in the Spurious Correlations: Boosting Moment Retrieval with Dynamic Learning, accessed August 28, 2025, <https://arxiv.org/html/2501.07305v2>