

Generative Localization: A Comprehensive Analysis of Emerging Paradigms in Object Detection and Video Moment Retrieval

Foundational Paradigms: Generative vs. Discriminative Localization

The field of computer vision, particularly in tasks requiring the localization of objects in space or time, has long been dominated by two fundamental modeling philosophies: discriminative and generative. While both fall under the umbrella of supervised learning, their underlying probabilistic principles and resulting architectural designs are profoundly different. Understanding this distinction is crucial for appreciating the recent paradigm shift towards generative localization, a movement that is reshaping the landscape of object detection and video moment retrieval. This section establishes the theoretical bedrock for this analysis, exploring the probabilistic foundations, core principles, and comparative merits of each approach.

Probabilistic Underpinnings of Localization Models

At the heart of the distinction between generative and discriminative models lies the probabilistic quantity they aim to model.¹ Given an input, such as an image represented by a feature vector

x , and a corresponding label, such as an object class and its bounding box coordinates represented by y , the two paradigms approach the modeling task from different perspectives.

A **discriminative model** focuses on learning the conditional probability distribution

$p(y|x)$ directly.⁴ This is the most natural and direct formulation for a classification or regression task. The model's sole purpose is to learn a decision boundary that separates different classes or to predict a continuous value for

y given x . It does not concern itself with how the data x is distributed or generated; it only cares about the boundary between classes.¹ For object detection, this means the model learns to directly map image pixels to a set of labels and bounding boxes. Prominent examples of discriminative models include Logistic Regression, Support Vector Machines (SVMs), and most conventional deep neural networks used for classification.⁴

Conversely, a **generative model** learns the joint probability distribution $p(x,y)$.¹ This is a more ambitious and holistic modeling goal. By capturing the joint distribution, the model learns the underlying patterns of the data and how it was generated.⁴ Typically, this is achieved by modeling the class-conditional probability

$p(x|y)$ and the prior probability of the class $p(y)$. Once these are learned, the desired conditional distribution for classification, $p(y|x)$, can be derived using Bayes' rule:

$$p(y|x) = \frac{p(x|y)p(y)}{p(x)}$$

The critical distinction is that modeling the joint distribution $p(x,y)$ allows the model to do more than just discriminate. It can be used to *generate* new, synthetic data samples (x,y) that are statistically consistent with the training data distribution.⁴ This generative capability is the defining feature and namesake of this class of models, which includes Naive Bayes, Hidden Markov Models, Generative Adversarial Networks (GANs), and Variational Autoencoders (VAEs).⁴ This fundamental difference in the probabilistic objective function leads to distinct architectural philosophies, performance characteristics, and application domains for each paradigm.

Core Principles of Discriminative Localization: Learning the Boundary

The history of modern object detection is largely the history of discriminative localization. This paradigm is characterized by its direct, predictive focus and its reliance on meticulously engineered pipelines to manage the complexity of the task. The core principle is to learn a mapping that discriminates objects from the background and from each other within a predefined set of categories.

The dominant approach has been the "propose-then-classify/regress" pipeline,

exemplified by the R-CNN family of models, including Faster R-CNN.⁷ This two-stage process first employs a Region Proposal Network (RPN) to generate a dense set of potential object locations (region proposals or anchors) across the image. In the second stage, a classification and regression head processes each proposal to determine if it contains an object and, if so, refines its bounding box coordinates.⁸ Single-stage detectors, such as the You Only Look Once (YOLO) family and Single Shot Detectors (SSD), streamline this process by performing classification and regression directly on a dense grid of predefined anchor boxes in a single forward pass.⁷

Regardless of the specific architecture, these methods share common characteristics rooted in their discriminative nature. They are built around hand-designed components that encode strong prior knowledge about the task. **Anchor boxes**, for instance, are predefined boxes of various sizes and aspect ratios that serve as initial guesses, assuming that objects will conform to these common shapes.¹³ Because these models generate a massive number of redundant and overlapping predictions (often thousands per image), they require a crucial post-processing step called

Non-Maximum Suppression (NMS). NMS is a heuristic algorithm that filters out duplicate detections by keeping only the box with the highest confidence score in a cluster of overlapping boxes.¹³

The primary strengths of the discriminative paradigm are its exceptional predictive performance and computational efficiency at inference time, especially for well-defined, closed-set tasks.¹ Models like YOLOv8 are renowned for their real-time speed, making them suitable for applications like autonomous driving and live video analytics.⁷ They also tend to converge faster during training compared to their generative counterparts.¹

However, this paradigm is not without significant limitations. Its heavy reliance on hand-designed components like anchors and NMS makes the pipeline complex and difficult to optimize end-to-end.¹³ The models are inherently limited in scope; they are trained to predict from a fixed set of categories and struggle to generalize to novel, unseen objects (an open-set problem).⁵ Furthermore, because they only learn the features necessary to form a decision boundary, they can be brittle and susceptible to out-of-distribution (OOD) data, where shifts in the input data distribution can cause a significant drop in performance.¹⁹ They can also inherit and amplify biases present in the training data, leading to skewed predictions.⁵

Core Principles of Generative Localization: Learning the Data Distribution

The generative approach to localization represents a fundamental shift in problem formulation. Instead of learning a discriminative boundary, the goal is to learn the underlying distribution of the data itself—to model *how* plausible scenes containing objects are generated.⁴ Localization is reframed as a generation task: given some form of input, the model's objective is to

generate the complete and correct set of localizations (e.g., bounding boxes or temporal spans). This reframing moves the field away from an indirect, filtering-based methodology toward a direct, creation-based one. While discriminative models answer thousands of simple questions like "is there an object in this specific box?", generative models directly address the holistic question: "what is the complete set of objects in this image?".

This generative philosophy has manifested in two primary architectural patterns for localization:

1. **Direct Set Prediction:** Pioneered by the DETection TRansformer (DETR), this approach formulates object detection as a direct set prediction problem.¹³ The model takes an image and a small, fixed set of learnable "object queries" as input and outputs the final set of predictions in a single, parallel pass. This end-to-end philosophy forces the model to learn the global relationships between objects and the overall image context, internally reasoning about object co-occurrence and avoiding duplicates. This structure entirely obviates the need for components like RPNs, anchors, and NMS, leading to a significantly simplified and more elegant pipeline.¹³
2. **Iterative Refinement (Denoising):** This approach, most prominently realized in diffusion models like DiffusionDet for object detection and MomentDiff for video moment retrieval, casts localization as a denoising process.²⁰ The model is trained to reverse a process that gradually adds noise to ground-truth localizations. At inference time, the model starts with a set of purely random localizations (e.g., noisy boxes sampled from a Gaussian distribution) and iteratively refines them over several steps, guided by the input image or text query, until they converge to the correct outputs.²⁰ This "noise-to-box" or "random-to-real" paradigm is inherently generative and also eliminates the need for predefined anchors or proposals.

The principal advantage of generative localization is its architectural simplicity and flexibility. By learning to generate the output set directly, these models internalize the logic that was previously handled by complex, hand-designed heuristics. This leads to fully end-to-end trainable systems. Furthermore, because generative models learn the entire data distribution, they possess several ancillary benefits. They can be used for data augmentation by generating new training samples, can handle missing or partially labeled data more gracefully, and are better equipped to model compositional scenes (e.g., a person wearing glasses and a hat) where discriminative models might need to see all possible combinations during training.¹

A Comparative Analysis of Strengths, Weaknesses, and Ideal Use Cases

The choice between a discriminative and a generative approach to localization is ultimately dependent on the specific requirements of the task. Each paradigm presents a distinct set of trade-offs.

Discriminative Models are the established workhorses of computer vision, optimized for speed and accuracy in closed-world scenarios.

- **Strengths:**
 - **High Predictive Accuracy:** Often achieve state-of-the-art accuracy on specific prediction tasks with fixed categories.⁵
 - **Inference Efficiency:** Typically faster at inference due to their single-pass, feed-forward nature, making them ideal for real-time applications.¹
 - **Faster Convergence:** Generally require fewer training epochs to converge compared to early generative models.¹
- **Weaknesses:**
 - **Pipeline Complexity:** Rely on complex, multi-stage pipelines with hand-designed components like anchors and NMS that are difficult to optimize jointly.¹³
 - **Limited Scope:** Primarily focused on prediction and lack the versatility to generate new data or handle tasks outside their training distribution.⁵
 - **Brittleness to OOD:** Performance can degrade significantly when faced with data from a different distribution than the training set.¹⁹
- **Ideal Use Cases:** Real-time object detection in controlled environments (e.g., autonomous driving, surveillance), medical imaging where object classes are well-defined, and any application where speed and accuracy on a fixed set of

categories are paramount.

Generative Models represent a more flexible and generalizable paradigm, trading some of the raw efficiency of discriminative models for a deeper understanding of the data.

- **Strengths:**

- **Architectural Simplicity:** Offer an end-to-end solution that eliminates the need for complex, hand-designed heuristics like NMS and anchors.¹³
- **Flexibility and Generalization:** Can handle open-set and open-vocabulary tasks more naturally, generate new data for augmentation, and are more robust to missing data and compositional variations.¹
- **Improved Robustness:** By learning the full data distribution, they are forced to build a richer world model. This acts as a form of implicit regularization, potentially leading to better generalization and robustness against OOD inputs. The model must learn not just the decision boundary but also the appearance of the data itself ($p(x|y)$), making it less susceptible to spurious correlations in the training set.¹

- **Weaknesses:**

- **Computational Cost:** Often require larger datasets and more computational resources for training.⁵ Early versions suffered from slow convergence.¹⁴
- **Inference Latency:** Iterative models like diffusion-based detectors can be slower at inference than single-pass discriminative models.¹
- **Interpretability:** Understanding how a generative model produces its output can be more challenging than interpreting a discriminative classifier.⁵

- **Ideal Use Cases:** Open-ended and open-vocabulary detection where the set of objects is not known beforehand, anomaly detection, data augmentation, and tasks requiring a deep contextual understanding of the scene.

In conclusion, while discriminative models have long been the standard for their raw predictive power, the field is increasingly moving towards generative approaches. The allure of a simplified, end-to-end pipeline combined with greater flexibility and robustness presents a compelling path forward, promising models that not only perceive the world but also possess a more fundamental understanding of its underlying structure.

The Generative Shift in Object Detection

The domain of object detection has been a primary battleground for the paradigm shift from discriminative to generative localization. For years, progress was measured by incremental improvements to complex, multi-stage pipelines. The introduction of generative concepts, first through Transformer-based set prediction and later through diffusion-based box generation, has fundamentally altered the field's trajectory. These new architectures have not only simplified the detection pipeline but have also unlocked new capabilities, such as open-ended detection, pushing the boundaries of what is possible.

Transformer-Based Set Prediction: The DETR Family and the Bipartite Matching Loss

The DETection TRansformer (DETR), introduced in 2020, marked the first major departure from the dense-prediction paradigm that had dominated object detection.¹³ It reframed the task not as a massive classification and regression problem over thousands of anchors, but as a direct set prediction problem: given an image, directly output the small, finite set of objects it contains. This conceptual shift was enabled by a novel architecture and a unique loss function.

Architectural Overview

The DETR architecture is conceptually simple and consists of three main components¹¹:

1. **CNN Backbone:** A standard convolutional neural network (e.g., ResNet) is used to extract a compact feature representation from the input image. This feature map is then flattened and supplemented with positional encodings to retain spatial information, which is crucial as the subsequent Transformer is permutation-invariant.²⁷
2. **Transformer Encoder-Decoder:** This is the core of the model.
 - The **Encoder** takes the sequence of image features as input. Its multi-head self-attention layers allow it to reason about the global context of the image,

modeling the relationships between different parts of the scene.²⁷

- The **Decoder** takes a small, fixed number (e.g., $N=100$) of learned positional embeddings, known as **object queries**, as input.²⁷ These queries are updated through self-attention and cross-attention with the encoder's output, allowing each query to specialize and focus on a specific object or region in the image. The decoding is done in parallel for all queries, a key difference from autoregressive language models.²⁸
- 3. **Feed-Forward Networks (FFNs):** Two prediction heads, implemented as FFNs, are attached to the output of the decoder. For each of the N output queries, one head predicts the class label (including a "no object" class), and the other predicts the bounding box coordinates (center x , center y , width, height).²⁷

The Bipartite Matching Loss

The most critical innovation of DETR is its loss function, which enables end-to-end training without requiring NMS.¹³ The loss computation involves two steps: first, finding an optimal one-to-one matching between the model's predictions and the ground-truth objects, and second, computing a loss based on this matching.

1. **Set Prediction and Padding:** The model outputs a fixed-size set of N predictions, $\{y^i\}_{i=1}^N$. The ground-truth set of objects, y , is variable in size. To create two sets of equal cardinality for matching, the ground-truth set is padded with "no object" ($\emptyset$) placeholders up to size N .²⁸
2. **Optimal Bipartite Matching:** The model must find the best possible one-to-one assignment between the predicted and ground-truth sets. This is framed as an assignment problem and solved efficiently using the **Hungarian algorithm**.²⁴ The algorithm finds a permutation of N elements, $\sigma \in S_N$, that minimizes the total matching cost:

$$\sigma^* = \arg \min_{\sigma \in S_N} \sum_{i=1}^N L_{\text{match}}(y_i, y^{\sigma(i)})$$
 Here, L_{match} is the pairwise matching cost between a ground-truth object y_i and a prediction at index $\sigma(i)$.
3. **Matching Cost Formulation:** The cost function is designed to account for both class and box similarity. For a ground-truth object $y_i = (c_i, b_i)$, where c_i is the class label and b_i is the bounding box, and a prediction $y^{\sigma(i)} = (p^{\sigma(i)}(c_i), b^{\sigma(i)})$, the matching cost is defined as²⁴:

$$L_{\text{match}}(y_i, y^{\sigma(i)}) = -1\{c_i \neq \emptyset\} p^{\sigma(i)}(c_i) + 1\{c_i \neq \emptyset\} L_{\text{box}}(b_i, b^{\sigma(i)})$$
 The first term is the negative log-probability of the correct class, making the cost

lower for confident, correct class predictions. The second term, L_{box} , is a bounding box loss, which is a linear combination of the L1 loss and the Generalized IoU (GIoU) loss.¹² This cost is only computed for actual objects, not the $\emptyset$ placeholders.

4. **Final Hungarian Loss:** After finding the optimal assignment σ^* , the final loss for the model is computed as a linear combination of the negative log-likelihood for class prediction (a form of cross-entropy) and the bounding box loss over all matched pairs:

$$L_{Hungarian}(y, y^*) = \frac{1}{N} \sum_i [-\log p(\sigma^*(i) | c_i) + 1_{\{c_i \neq \emptyset\}} L_{box}(b_i, b^{\sigma^*(i)})]$$

This set-based loss forces each prediction to be unique, as the model is heavily penalized for producing duplicate detections for the same object, thereby learning an implicit NMS-like behavior.¹³

Evolution and Variants

While revolutionary, the original DETR had notable drawbacks, including extremely slow training convergence (requiring up to 500 epochs) and inferior performance on small objects compared to established detectors like Faster R-CNN.¹⁴ These issues spurred a wave of research leading to significant improvements.

Deformable DETR addressed both problems by introducing a more efficient, deformable attention module that only attends to a small, learned set of sampling points around a reference point, rather than all pixels in the feature map. This reduced computational complexity and allowed for the processing of multi-scale feature maps, dramatically improving small object detection and accelerating convergence by more than 10x.³³ More recently, models like

RT-DETR have further optimized the architecture, particularly the decoder, to achieve real-time performance, making them directly competitive with, and in some cases superior to, the highly optimized YOLO family.¹⁵

Denoising Diffusion for Box Generation: The Architecture and Methodology of DiffusionDet

A second, distinct generative paradigm for object detection emerged with the application of denoising diffusion models. **DiffusionDet** formulates object detection not as a single-pass set prediction, but as an iterative refinement process that generates object boxes from random noise.²⁰ This "noise-to-box" philosophy provides a powerful and flexible alternative to both dense-prediction and query-based methods.

Conceptual Framework

The core concept of DiffusionDet is to treat object detection as a generative task over the space of bounding boxes.²⁰ The model learns to reverse a Markovian diffusion process that systematically corrupts ground-truth bounding boxes by adding Gaussian noise over a series of timesteps. During inference, the model starts with a set of boxes sampled from a pure Gaussian distribution and applies the learned reverse process to iteratively denoise them, eventually yielding the final, accurate object detections.²⁰ This approach is appealing because it starts from a simple, parameter-free random state, eliminating the need for either hand-designed anchors or learnable queries.²⁰

Architecture

The DiffusionDet architecture is cleanly separated into two main components³⁶:

1. **Image Encoder:** Similar to DETR, a standard CNN or Transformer backbone (e.g., ResNet with an FPN) is used to process the input image once and extract multi-scale feature maps. This deep feature representation serves as the condition for the denoising process.²⁰
2. **Detection Decoder:** This module is responsible for the iterative refinement of the bounding boxes. Its design is heavily inspired by Sparse R-CNN and consists of **6 cascading stages**.²⁰ In each stage, the current set of proposal boxes is used to crop RoI features from the image feature maps. These features are then fed into a detection head, which performs classification and regresses a refined set of boxes. The output boxes of one stage become the input proposals for the next,

allowing for progressive refinement. A key distinction from Sparse R-CNN is that DiffusionDet's decoder is designed to be reused iteratively across multiple denoising steps, with its behavior modulated by a timestep embedding that signals the current noise level.²⁰

Training and Inference Process

The training and inference stages in DiffusionDet correspond to the reverse and forward diffusion processes, respectively.

- **Training (Learning the Reverse Process):**

1. **Box Corruption:** For a given image, the set of ground-truth boxes is padded with random boxes to a fixed number, N_{train} .³⁶
2. A random timestep t is sampled, and a corresponding amount of Gaussian noise is added to the ground-truth boxes, creating a set of noisy boxes b_t .²⁰
3. These noisy boxes b_t are fed into the detection decoder, which is conditioned on the image features. The model is trained to predict the original, noise-free ground-truth boxes b_0 .²⁰
4. The training objective is a **set prediction loss**, similar to that used in DETR and Sparse R-CNN, which is applied between the decoder's predictions and the ground-truth boxes. An optimal transport assignment method is used to match multiple predictions to each ground truth, and the loss combines classification and box regression terms.²⁰

- **Inference (The Denoising Process):**

1. **Initialization:** The process begins by sampling an initial set of Neval boxes from a standard Gaussian distribution.²⁰
2. **Iterative Refinement:** The model then iteratively denoises this set of boxes. In each step, the current set of noisy boxes is passed through the detection decoder to predict a clean version of the boxes.
3. **Sampling:** A diffusion sampler, such as DDIM (Denoising Diffusion Implicit Models), uses this prediction to compute the estimated boxes for the next, slightly less noisy, timestep.²⁰
4. This process is repeated for a predefined number of steps, progressively refining the random boxes into the final set of accurate predictions. A key advantage of DiffusionDet is its flexibility: the number of boxes used during inference (Neval) can be different from the number used during training (N_{train}), and the number of iterative refinement steps can be adjusted to

trade off accuracy for speed, all without retraining the model.²⁰

Beyond Predefined Categories: Open-Ended Detection with Generative Language

The generative paradigm has enabled an even more ambitious vision for object detection: moving beyond any predefined set of categories to a truly open-ended system. While open-vocabulary object detection (OVOD) was a significant step, allowing models to detect objects described by arbitrary text prompts provided at inference time, it still required a user to supply the potential category names.²² The latest generative models eliminate this final constraint.

Problem Redefinition: Generative Open-Ended Object Detection

This new setting, termed **generative open-ended object detection**, formulates the task as a joint localization and free-form text generation problem.²² The model is not given a list of categories to match against; instead, it must both locate all objects of interest in an image and

generate their names from its own internal vocabulary.²² This is a more general and practical scenario, as it does not assume the user has prior knowledge of the objects present in an image and avoids the ambiguity of manually defining category names (e.g., "person" vs. "a man" vs. "young boy").²² This represents a profound unification of perception and language, where detection becomes a task of grounded language generation.

The GenerateU Framework

The pioneering work in this area is the **GenerateU** framework, which seamlessly integrates a visual detector and a language model in an end-to-end architecture²²:

1. **Visual Object Detector:** A class-agnostic object detector, specifically

Deformable DETR, is used as a region proposal generator. Its role is to identify and localize all potential object regions in the image without assigning them to a specific class.²²

2. **Language Model:** A multimodal large language model (MLLM) with an encoder-decoder architecture takes the visual features corresponding to each detected region as input. The decoder then autoregressively generates a free-form text string that names or describes the object in that region.²²

The entire model is trained end-to-end. To broaden its knowledge base beyond limited annotated detection datasets, GenerateU is trained on a combination of human-annotated object-language paired data (like Visual Genome) and massive-scale image-text web data (like GrIT-20M).²² This strategy allows the language model to learn a vast vocabulary and associate it with visual concepts, enabling it to generate accurate and descriptive names for a wide array of objects. A region-word alignment loss and a pseudo-labeling strategy are employed to provide additional supervision and further enhance the model's capabilities.³⁸

Comparative Analysis and Performance on the COCO Benchmark

The practical viability of these generative models is best assessed through their performance on standard benchmarks like COCO. A comparative analysis reveals that generative approaches have not only become competitive but are now setting the state of the art.

Model Category	Model	Backbone	Training	AP	AP ₅₀	AP ₇₅	AP _S	AP _M	AP _L	FPS (T4 GPU)
Discriminative	Faster R-CNN ²⁰	ResNet-50	-	40.2	-	-	-	-	-	~18 ₁₀
	YOLOv8-X ¹⁵	-	-	53.9	71.0	58.7	35.7	59.3	70.7	50

Gen erati ve	DETR ²⁰	ResNet-50	500 epochs	42.0	-	-	-	-	-	~28 ¹³
	Deformable DETR ³³	ResNet-50	50 epochs	45.7	64.7	49.4	28.5	49.0	60.1	~28
	RT-DETR-R50 ¹⁵	ResNet-50	-	53.1	-	-	34.5	55.9	68.3	108
	RF-DETR-M ⁴⁰	-	-	54.7	73.6	-	-	-	-	221
	Diffusion Det ²⁰	ResNet-50	1 step	45.8	-	-	-	-	-	-
	Diffusion Det ⁴¹	ResNet-50	4 steps	46.1	-	-	-	-	-	-

Table 1: Comparative Performance of Object Detection Models on the COCO Benchmark. This table contrasts leading discriminative models (Faster R-CNN, YOLOv8) with key generative models (DETR variants, DiffusionDet). AP values are reported on the COCO val2017 set. FPS values are approximate and depend on implementation and hardware.

The data reveals a clear trajectory. The original DETR was competitive with Faster R-CNN but slower and worse on small objects.¹³ However, subsequent innovations have dramatically changed the landscape. RT-DETR, a real-time variant, now outperforms the state-of-the-art YOLOv8-L model in both speed and accuracy, achieving 53.1 AP at 108 FPS versus YOLOv8-L's 52.9 AP at 71 FPS.¹⁵ RF-DETR further pushes performance, with its medium version reaching 54.7 AP at an even higher framerate.⁴⁰ DiffusionDet also demonstrates strong performance, outperforming both Faster R-CNN and the original DETR with just a single inference step.²⁰

Regarding robustness, benchmarks on OOD datasets like COCO-O have shown that performance is heavily influenced by the backbone and pre-training data. While early Transformer-based detectors did not inherently offer more robustness, the integration of large-scale foundation models as backbones has led to a significant leap in OOD generalization for all architectures.¹⁹ This suggests that the generative architecture, when paired with powerful, broadly trained feature extractors, can achieve superior robustness.

Generative Approaches to Video Moment Retrieval

The paradigm shift from discriminative to generative modeling is not confined to the spatial domain of object detection; it is also making significant inroads into the temporal domain of Video Moment Retrieval (VMR). VMR aims to localize a specific temporal segment within an untrimmed video that corresponds to a natural language query. Traditional methods have mirrored the dense-proposal approach of early object detectors, but recent generative models are introducing a more efficient, robust, and intuitive framework for temporal localization.

From Proposal Ranking to Direct Generation: A New VMR Paradigm

The conventional approach to VMR has been predominantly discriminative and can be categorized as a "proposal-rank" pipeline.²¹ In this two-stage process, a model first generates a large number of candidate temporal segments, or "proposals." These proposals can be created using simple sliding windows of various scales, predefined temporal anchors, or more sophisticated methods like constructing a 2D temporal map where each point represents a potential start-end pair, as seen in the 2D-TAN model.²¹ In the second stage, each of these proposals is scored based on its semantic similarity to the given text query, and the highest-scoring proposal is selected as the final prediction.⁴³

This methodology, while effective, suffers from several inherent drawbacks. The generation of dense proposals is computationally expensive and creates significant redundancy, with the vast majority of proposals being irrelevant negative samples.²¹

This extreme imbalance between positive and negative samples can hinder training efficiency.⁴³ More critically, these models are prone to learning and exploiting spurious temporal biases present in the training datasets. For instance, if most annotated moments in a dataset tend to occur in the middle of videos, the model may develop a strong prior to favor proposals from that region, regardless of the semantic content, thereby limiting its generalization to real-world scenarios where events can occur at any time.²¹

Diffusion Models for Temporal Localization: The MomentDiff Framework

In response to the limitations of discriminative methods, **MomentDiff** was introduced as the first framework to tackle VMR from a purely generative perspective.²¹ It recasts the task as a denoising diffusion process, which is conceptually analogous to a human quickly scanning a video: starting from a random guess and gradually zeroing in on the correct moment. This "random-to-real" approach fundamentally changes the localization process.²¹

Architectural Breakdown

The MomentDiff framework is elegantly designed with two primary components that work in concert to achieve generative localization²¹:

1. **Similarity-aware Condition Generator (SCG):** This module is responsible for creating a rich, multimodal representation that guides the denoising process. It takes pre-extracted video and text features as input and uses a series of cross-attention and self-attention layers to fuse them. The output is a "fusion embedding" that encodes the frame-by-frame semantic similarity between the video content and the text query. This embedding acts as the conditional information for the diffusion model, providing the necessary semantic cues to guide the localization.
2. **Video Moment Denoiser (VMD):** This is the core generative engine of the framework. It is a Transformer-based network that implements the reverse diffusion process. The VMD takes two inputs: a noisy temporal span, represented by its normalized center time and duration length, and the conditional fusion

embedding from the SCG. Its task is to predict a less noisy, more accurate version of the temporal span.

Loss Functions and Training

The model is trained end-to-end using a composite loss function that optimizes both modules simultaneously ²¹:

- Similarity Loss (L_{sim}):** This loss trains the SCG. It is a combination of a pointwise cross-entropy loss, which encourages high similarity scores for frames within the ground-truth moment, and a pairwise margin loss, which ensures that positive frames have higher scores than negative frames. The full formulation is:

$$L_{sim} = -N \sum_i [y_i \log(s_i) + (1 - y_i) \log(1 - s_i)] + N \sum_j \max[0, \beta + s_{nj} - s_{pj}]$$
 where s_i is the similarity score for frame i , y_i is the binary ground-truth label, and s_{pj} and s_{nj} are scores for sampled positive and negative frames, respectively.²¹
- Denoising Loss (L_{vmr}):** This loss trains the VMD to perform accurate localization. It is a weighted sum of three components: an L1 loss for the raw coordinates, a generalized IoU loss to measure the overlap between the predicted and ground-truth spans, and a cross-entropy loss for a predicted confidence score. The formulation is:

$$L_{vmr}(x_0, f_\theta(x_m, m, F)) = \lambda L1 // x_0 - x^{m-1} // + \lambda_{iou} Liou(x_0, x^{m-1}) + \lambda_{ce} Lce(z^{m-1})$$
 where x_0 is the ground-truth span and x^{m-1} and z^{m-1} are the predicted span and confidence score from the denoiser f_θ .²¹

Inference Process

During inference, MomentDiff operationalizes its generative nature. It begins by sampling a set of temporal spans from a standard Gaussian distribution. These completely random spans are then fed into the VMD, along with the SCG's fusion embedding for the given video-query pair. The VMD predicts a cleaner set of spans, and this process is repeated for a fixed number of steps, with each step producing a more refined estimate. The final prediction is the span with the highest confidence score after the last denoising step.²¹ Because this process can start from any random point in the temporal space and is guided by semantics rather than a learned location

prior, it is inherently more robust to the temporal location biases that plague discriminative models.²¹

The Role of Multimodal Large Language Models in Open-Ended Retrieval

The most recent evolution in VMR leverages the extraordinary capabilities of Multimodal Large Language Models (MLLMs), pushing the task even further into the generative domain. Models like **Mr. BLIP** represent a paradigm shift by casting VMR as an open-ended **sequence-to-sequence problem**.⁴⁷

In this framework, a powerful, pre-trained MLLM (like BLIP-2) is adapted to process video inputs. The model is given a sequence containing the text query interleaved with visual tokens from sampled video frames and their corresponding timestamps. Instead of having a specialized prediction head that outputs numerical coordinates, the MLLM is trained to generate the answer directly as a text string, such as: "The moment where 'a person is petting a dog' occurs from to".⁴⁷

This approach offers two significant advantages. First, it completely removes the constraint of predicting a fixed number of moments. The model can generate text describing one, two, or any number of relevant moments that match the query, reflecting the true nature of the video's content.⁴⁷ Second, it harnesses the immense contextual and temporal reasoning abilities that MLLMs acquire during their extensive pre-training on web-scale data, leading to a more nuanced understanding of complex queries and scenes.⁴⁷

Performance Evaluation on Charades-STA and ActivityNet Captions

The efficacy of the generative paradigm in VMR is substantiated by strong empirical results on widely used benchmarks. MomentDiff, in particular, has demonstrated state-of-the-art performance, outperforming previous discriminative and DETR-based methods.

Dataset	Method	Feature	R1@IoU=0.5	R1@IoU=0.7	mAP@avg
---------	--------	---------	------------	------------	---------

		Backbone			
Charades-STA	2D-TAN★ ²¹	VGG	41.34	23.91	29.26
	VSLNet★ ²¹	C3D	48.67	30.33	30.16
	MomentDETR★ ²¹	SF+C	53.22	30.87	30.43
	MomentDiff ₂₁	VGG	51.94	28.25	31.66
	MomentDiff ₂₁	C3D	53.79	30.18	31.89
	MomentDiff ₂₁	SF+C	55.57	32.42	32.85
ActivityNet Captions	2D-TAN ²¹	C3D	44.63	27.53	27.26
	MMN ²¹	C3D	48.26	28.95	28.74
	MomentDETR ²¹	C3D	43.19	25.74	25.63
	MomentDiff ₂₁	C3D	46.52	28.43	28.19

Table 2: Comparative Performance of Video Moment Retrieval Models. This table compares the performance of the generative MomentDiff model against leading discriminative methods (2D-TAN, VSLNet, MMN) and the Transformer-based MomentDETR. Metrics are "Recall@1" at different IoU thresholds and mean Average Precision. Results are from the Charades-STA and ActivityNet Captions datasets. '★' indicates re-implementation under a consistent scheme.

As shown in Table 2, MomentDiff consistently outperforms prior methods across all feature backbones on the Charades-STA dataset. With powerful SF+C features, it achieves an R1@0.5 of 55.57, a significant improvement over MomentDETR's 53.22.²¹ On the ActivityNet Captions dataset, it also surpasses MomentDETR, achieving an R1@0.5 of 46.52 compared to 43.19, while being orders of magnitude more efficient in

training and inference time than older proposal-based methods like 2D-TAN and MMN.²¹

Crucially, the generative nature of MomentDiff provides benefits beyond standard accuracy metrics. By training the model to denoise spans from arbitrary random locations, it is forced to rely more on the semantic guidance from the text-video fusion rather than on spurious location priors learned from the dataset. This makes it inherently more robust to shifts in the temporal distribution of events. This was validated on custom "anti-bias" datasets, where MomentDiff exhibited superior generalization and robustness compared to its discriminative counterparts, confirming that the generative approach is a powerful tool for mitigating the effects of dataset bias.²¹

Overarching Challenges and Technical Hurdles in Generative Localization

Despite the significant conceptual advancements and impressive empirical results, the shift towards generative localization is not without its challenges. These models introduce a new set of technical hurdles related to computational cost, performance on fine-grained tasks, and issues of controllability and bias. Addressing these challenges is the primary focus of current research and is essential for the widespread, practical deployment of generative localization systems in both object detection and video moment retrieval.

Computational Complexity, Training Stability, and Inference Latency

A significant barrier to the adoption of generative models has been their substantial computational demand compared to their discriminative predecessors. This manifests in several ways:

- **Training Cost and Convergence:** Generative models, especially those based on the Transformer architecture or diffusion processes, are notoriously data-hungry and require extensive training schedules to converge effectively.⁵ The original DETR model, for instance, required 500 epochs of training on the COCO dataset

to reach its peak performance, a stark contrast to the 12-50 epochs typically needed for models like Faster R-CNN.¹⁴ This prolonged training not only increases computational cost but also slows down the research and development cycle. Diffusion models add another layer of complexity, as their training can be unstable and highly sensitive to the choice of noise schedule and other hyperparameters.²⁵

- **Inference Latency:** While training cost is a concern for developers, inference latency is critical for real-world applications. The iterative nature of diffusion models presents a direct trade-off between accuracy and speed. A model like DiffusionDet can achieve respectable performance with a single denoising step, making it relatively fast. However, its performance consistently improves with more steps (e.g., from 45.5 AP with 1 step to 46.1 AP with 4 steps on COCO).⁴¹ This multi-step refinement process is inherently slower than the single forward pass of a highly optimized discriminative model like YOLO.²⁰ This makes deploying diffusion-based detectors in real-time scenarios, such as autonomous vehicles, a significant challenge that requires further algorithmic and hardware optimization.

The Persistent Challenge of Fine-Grained and Small Object Localization

Another well-documented limitation of early generative models is their difficulty with fine-grained localization, particularly in detecting small objects or precisely identifying the boundaries of short temporal moments.

- **Small Object Detection:** The original DETR architecture struggled to detect small objects effectively, showing a significant performance gap compared to detectors equipped with Feature Pyramid Networks (FPNs).²⁶ This was largely attributed to the Transformer encoder processing features at a single, often coarse, spatial resolution. While subsequent models like Deformable DETR have mitigated this by incorporating multi-scale features, the challenge remains an active area of research. For diffusion models, generating small objects presents a unique set of difficulties. The limited spatial information and subtle visual cues of small objects can be easily lost or overwhelmed by noise during the diffusion process, making them difficult to recover accurately.⁵⁰ Aligning cross-modal attention between a text prompt and a small visual region is also particularly challenging for generative models.⁵²
- **Fine-Grained VMR:** In video moment retrieval, this challenge manifests as difficulty in precisely localizing very short moments.⁵³ The subtle differences that

define the start and end boundaries of a brief event can be hard to capture and generate from a noisy prior. Discriminative models that explicitly regress boundaries or rank dense, fine-grained proposals may have an advantage in these scenarios. The problem is compounded by the fact that fine-grained visual categorization itself is difficult, as the differences between sub-categories can be extremely subtle and easily masked by variations in pose, lighting, and background clutter.⁵⁴

These issues point to a "generative trilemma" at the core of current research: a difficult trade-off between **inference speed**, **generative purity** (reliance on minimal hand-designed priors), and **fine-grained accuracy**. Discriminative models like YOLO excel at speed and fine-grained accuracy but sacrifice generative purity by relying on anchors and NMS. The original DETR prioritized purity but at the cost of speed and small-object accuracy. DiffusionDet also prioritizes purity but creates a direct trade-off between accuracy and inference speed. No single model has yet optimally solved all three challenges, indicating this is a fundamental frontier for future work.

Controllability, Interpretability, and Mitigating Inherent Model Biases

As generative models become more powerful, ensuring their outputs are reliable, understandable, and fair becomes increasingly critical.

- **Controllability:** The stochastic nature of the generation process, especially in diffusion models, makes their outputs inherently variable and difficult to control with absolute precision.⁴⁹ Small, random variations in the initial noise or early denoising steps can propagate and amplify, leading to different final outputs. While this diversity can be a feature, it becomes a problem in applications requiring deterministic and guaranteed-safe behavior. Ensuring that a generative model will consistently adhere to specific constraints or avoid generating harmful or unsafe content is a significant challenge.⁴⁹
- **Interpretability:** The "black box" nature of deep learning models is a well-known issue, and it can be exacerbated in complex generative frameworks. While techniques like visualizing attention maps in Transformers can provide some insight into the model's decision-making process⁵, the intricate, multi-step dynamics of a diffusion model are far less transparent. Understanding precisely *why* a model generated a particular bounding box or hallucinated an object is a difficult task, making debugging and trust-building more challenging.⁵

- **Bias Amplification:** Generative models are trained on vast, often web-scraped, datasets that contain implicit societal biases, stereotypes, and harmful content. There is a significant risk that these models will not only learn and reproduce these biases but also amplify them, generating outputs that reinforce existing inequalities.⁴⁹ The ability of these models to create novel content makes the potential for generating biased or harmful material a serious ethical concern that requires robust mitigation strategies, such as careful data filtering, adversarial training, and the incorporation of fairness constraints into the model's objective function.⁴⁹

Future Trajectories and Concluding Remarks

The ongoing transition from discriminative to generative localization marks a pivotal moment in the evolution of computer vision. By reframing perception as a generative task, researchers are creating simpler, more flexible, and more powerful systems that move beyond mere pattern recognition toward a deeper, more holistic understanding of visual data. While significant challenges remain, the current trajectory points toward a future where unified, multimodal generative models become the backbone of intelligent perception systems, extending their capabilities into complex 3D and spatio-temporal domains.

The Path Towards Unified, Multimodal Generative Perception

The progression from DETR's set prediction to GenerateU's free-form language generation in object detection, and from proposal-ranking to sequence-to-sequence generation in video moment retrieval, signals a clear trend: the convergence of vision and language within a unified generative framework.²² The future of localization is unlikely to lie in specialized, single-task models. Instead, the field is moving towards large-scale, pre-trained foundation models that can perform a wide array of perception tasks by generating outputs in a common, often textual, format.⁵⁸

Imagine a single, powerful model that can receive an image or video and a text prompt specifying a task. If the prompt is "detect all objects," the model generates a list of

object names and their coordinates. If the prompt is "where is the person walking the dog?", it generates the relevant temporal span. If the prompt is "describe this scene," it generates a detailed caption. This vision of a general-purpose perception engine is made possible by the generative paradigm, which naturally accommodates diverse outputs and leverages the immense reasoning capabilities of large language models.⁵⁸

Extending Generative Localization to 3D and Complex Spatio-Temporal Scenarios

The principles of generative localization are not confined to 2D images. The next frontier is the application of these concepts to more complex and higher-dimensional data, which is critical for applications like robotics, augmented reality, and autonomous driving.

- **3D Object Detection:** The ability to accurately perceive and localize objects in 3D space is essential for autonomous navigation. While early methods relied heavily on expensive LiDAR sensors, there is a strong push towards image-based 3D detection. Generative models are poised to play a key role here. Models like Generative Adversarial Networks (GANs), Neural Radiance Fields (NeRFs), and diffusion models are being explored to generate and reconstruct 3D shapes and scenes from 2D images.⁶² By learning the underlying generative process of how 3D scenes are projected into 2D images, these models can infer depth and 3D bounding boxes, a task that is fundamentally ill-posed for purely discriminative approaches.⁶⁵
- **Complex Video Understanding:** The capabilities of VMR are expanding beyond simple event localization to the understanding of complex, long-term interactions between multiple agents and objects. This involves tasks like spatio-temporal action localization and the generation of dynamic scene graphs that represent who did what to whom, and when. Generative models are exceptionally well-suited for modeling the high-dimensional, structured distributions of these complex interactions over time, promising a richer and more detailed understanding of video content.⁶⁸

Concluding Analysis and Recommendations for Future Research

This report has charted the rise of generative localization, a paradigm shift that is fundamentally reshaping object detection and video moment retrieval. By moving from learning decision boundaries to modeling the underlying data distribution, generative approaches have enabled simpler, end-to-end architectures that are more flexible, robust, and generalizable than their discriminative predecessors. Models like DETR, DiffusionDet, and MomentDiff have not only achieved state-of-the-art performance on standard benchmarks but have also unlocked new frontiers in open-ended detection and robust temporal localization.

However, the journey is far from over. The full potential of generative localization will only be realized by addressing the critical challenges that currently limit its practical application. Based on this comprehensive analysis, the following areas are recommended as key priorities for future research:

1. **Efficiency and Scalability:** Developing more computationally efficient Transformer and diffusion architectures is paramount. Research into model quantization, knowledge distillation, and faster sampling techniques for diffusion models will be crucial for enabling real-time performance and deployment on resource-constrained devices.⁵⁰
2. **Fine-Grained Representation Learning:** A primary focus must be on enhancing the ability of generative models to capture and synthesize fine-grained details. This involves developing new architectural components and training strategies specifically designed to improve performance on small objects and short temporal moments, closing the gap with specialized discriminative methods.⁵¹
3. **Controllability, Safety, and Ethical Alignment:** As these models become more capable, research into making them more controllable, interpretable, and safe is essential. Developing novel methods for guiding the generative process, auditing models for fairness, and mitigating the amplification of societal biases must be a central part of the research agenda to ensure these technologies are deployed responsibly.⁴⁹
4. **Harnessing Generative Power for Data Scarcity:** The inherent ability of these models to generate synthetic data should be leveraged more systematically to address long-standing challenges in domains with scarce annotated data, such as rare disease detection in medical imaging or defect identification in industrial manufacturing.⁷²

In conclusion, the generative paradigm is not merely an alternative to discriminative methods; it represents a more fundamental and powerful approach to machine perception. By continuing to innovate in these key research areas, the computer vision community can build upon the remarkable progress of recent years and develop the

next generation of intelligent systems that can perceive and interact with the world in a truly comprehensive and human-like manner.

Works cited

1. Generative versus Discriminative Methods for Object ... - Microsoft, accessed August 15, 2025,
<https://www.microsoft.com/en-us/research/wp-content/uploads/2016/05/Bishop-CVPR-05.pdf>
2. Comparison of Generative and Discriminative Techniques for Object Detection and Classification - University of Edinburgh Research Explorer, accessed August 15, 2025,
<https://www.research.ed.ac.uk/en/publications/comparison-of-generative-and-discriminative-techniques-for-object>
3. Generative versus Discriminative Methods for Object Recognition - Microsoft Research, accessed August 15, 2025,
<https://www.microsoft.com/en-us/research/publication/generative-versus-discriminative-methods-object-recognition/>
4. What is the difference between a generative and a discriminative algorithm? [closed], accessed August 15, 2025,
<https://stackoverflow.com/questions/879432/what-is-the-difference-between-a-generative-and-a-discriminative-algorithm>
5. Generative Vs Discriminative: Understanding Machine Learning Models | by Kanerika Inc, accessed August 15, 2025,
<https://medium.com/@kanerika/generative-vs-discriminative-understanding-machine-learning-models-87e3d2b3b99f>
6. Understanding Generative and Discriminative Models | by BotPenguin - Chatbots Life, accessed August 15, 2025,
<https://blog.chatbotslife.com/understanding-generative-and-discriminative-models-a27dc91ebe1f>
7. YOLO vs. Faster R-CNN: Which Object Detection Framework Should You Use for Real-Time Tasks? : r/deeplearning - Reddit, accessed August 15, 2025,
https://www.reddit.com/r/deeplearning/comments/1luqx6z/yolo_vs_faster_rcnn_w_high_object_detection/
8. A Survey of Deep Learning-based Object Detection - ResearchGate, accessed August 15, 2025,
https://www.researchgate.net/publication/334623937_A_Survey_of_Deep_Learning-based_Object_Detection
9. Sparse R-CNN: End-to-End Object Detection With Learnable Proposals - CVF Open Access, accessed August 15, 2025,
https://openaccess.thecvf.com/content/CVPR2021/papers/Sun_Sparse_R-CNN_End-to-End_Object_Detection_With_Learnable_Proposals_CVPR_2021_paper.pdf
10. Comparative Analysis of YOLO and Faster R-CNN Models for Detecting Traffic Object - The Science and Information (SAI) Organization, accessed August 15, 2025,

https://thesai.org/Downloads/Volume16No3/Paper_42-Comparative_Analysis_of_YOLO_and_Faster_R_CNN_Models.pdf

11. DETection TRansformer (DETR) vs. YOLO for object detection. | by Fahim Rustamy, PhD, accessed August 15, 2025, <https://medium.com/@faheemrustamy/detection-transformer-detr-vs-yolo-for-object-detection-baeb3c50bc3>
12. Introduction to DETR (Detection Transformers): Everything You Need to Know - Lightly, accessed August 15, 2025, <https://www.lightly.ai/blog/detr>
13. End-to-End Object Detection with Transformers - Meta Research, accessed August 15, 2025, <https://research.facebook.com/publications/end-to-end-object-detection-with-transformers/>
14. detrex: Benchmarking Detection Transformers - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2306.07265>
15. DETRs Beat YOLOs on Real-time Object Detection - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2304.08069v3>
16. (PDF) Generative versus Discriminative Methods for Object Recognition - ResearchGate, accessed August 15, 2025, https://www.researchgate.net/publication/4156432_Generative_versus_discriminative_methods_for_object_recognition
17. YOLOv8 vs Faster R-CNN: A Comparative Analysis - Keylabs, accessed August 15, 2025, <https://keylabs.ai/blog/yolov8-vs-faster-r-cnn-a-comparative-analysis/>
18. Analysis of the performance of Faster R-CNN and YOLOv8 in detecting fishing vessels and fishes in real time - PubMed Central, accessed August 15, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11157610/>
19. COCO-O: A Benchmark for Object Detectors under Natural Distribution Shifts - CVF Open Access, accessed August 15, 2025, https://openaccess.thecvf.com/content/ICCV2023/papers/Mao_COCO-O_A_Benchmark_for_Object_Detectors_under_Natural_Distribution_Shifts_ICCV_2023_paper.pdf
20. DiffusionDet: Diffusion Model for Object Detection - CVF Open Access, accessed August 15, 2025, https://openaccess.thecvf.com/content/ICCV2023/papers/Chen_DiffusionDet_Diffusion_Model_for_Object_Detection_ICCV_2023_paper.pdf
21. MomentDiff: Generative Video Moment Retrieval from Random to Real, accessed August 15, 2025, <https://arxiv.org/pdf/2307.02869>
22. arxiv.org, accessed August 15, 2025, <https://arxiv.org/html/2403.10191v1>
23. Generative Region-Language Pretraining for Open-Ended Object Detection, accessed August 15, 2025, https://openaccess.thecvf.com/content/CVPR2024/papers/Lin_Generative_Region-Language_Pretraining_for_Open-Ended_Object_Detection_CVPR_2024_paper.pdf
24. DETR:End-to-End Object Detection with Transformers | Learning-Deep-Learning, accessed August 15, 2025, https://patrick-llgc.github.io/Learning-Deep-Learning/paper_notes/detr.html

25. What are common pitfalls encountered during diffusion model training? - Milvus, accessed August 15, 2025, <https://milvus.io/ai-quick-reference/what-are-common-pitfalls-encountered-during-diffusion-model-training>
26. For COCO validation see the comparison of DETR with Faster R-CNN with... - ResearchGate, accessed August 15, 2025, https://www.researchgate.net/figure/For-COCO-validation-ser-the-comparison-of-DETR-with-Faster-R-CNN-with-ResNet-50-and_fig40_360066821
27. What is Detection Transformers (DETR)? - Zilliz Learn, accessed August 15, 2025, <https://zilliz.com/learn/detection-transformers-detr-end-to-end-object-detection-with-transformers>
28. DETR - Hugging Face, accessed August 15, 2025, https://huggingface.co/docs/transformers/v4.40.1/en/model_doc/detr
29. A Review of DETECTION TRANSFORMER: From Basic Architecture to Advanced Developments and Visual Perception Applications - PubMed Central, accessed August 15, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12252279/>
30. DETR | Object Detection | Facebook AI | VisionWizard - Medium, accessed August 15, 2025, <https://medium.com/visionwizard/detr-b677c7016a47>
31. Introduction to DETR - Part 2: The Crucial Role of the Hungarian ..., accessed August 15, 2025, <https://www.digitalocean.com/community/tutorials/introduction-detr-hungarian-algorithm-2>
32. Benchmarking DETR vs. Faster R-CNN for Real-World Car Detection (KITTI) - Medium, accessed August 15, 2025, <https://medium.com/@monishatemp20/faster-r-cnn-vs-detr-a-practical-lightweight-benchmark-for-real-world-car-detection-on-kitti-983c714e2da3>
33. Deformable DETR: Deformable Transformers for End-to-End Object Detection. - GitHub, accessed August 15, 2025, <https://github.com/fundamentalvision/Deformable-DETR>
34. (PDF) DiffusionDet: Diffusion Model for Object Detection - ResearchGate, accessed August 15, 2025, https://www.researchgate.net/publication/365487783_DiffusionDet_Diffusion_Model_for_Object_Detection
35. [2211.09788] DiffusionDet: Diffusion Model for Object Detection - arXiv, accessed August 15, 2025, <https://arxiv.org/abs/2211.09788>
36. DiffusionDet Diffusion Model for Object Detection | Qiang Zhang, accessed August 15, 2025, <https://zhangtemplar.github.io/diffusion-det/>
37. [2403.10191] Generative Region-Language Pretraining for Open-Ended Object Detection - arXiv, accessed August 15, 2025, <https://arxiv.org/abs/2403.10191>
38. [Literature Review] Generative Region-Language Pretraining for ..., accessed August 15, 2025, <https://www.themoonlight.io/en/review/generative-region-language-pretraining-for-or-open-ended-object-detection>
39. FoundationVision/GenerateU: [CVPR2024] Generative Region-Language Pretraining for Open-Ended Object Detection - GitHub, accessed August 15,

- 2025, <https://github.com/FoundationVision/GenerateU>
40. Benchmarks - RF-DETR, accessed August 15, 2025, <https://rfdetr.roboflow.com/learn/benchmarks/>
 41. ShoufaChen/DiffusionDet: [ICCV2023 Best Paper Finalist] PyTorch implementation of DiffusionDet (<https://arxiv.org/abs/2211.09788>) - GitHub, accessed August 15, 2025, <https://github.com/ShoufaChen/DiffusionDet>
 42. COCO-O: A Benchmark for Object Detectors under Natural Distribution Shifts - ICCV 2023 Open Access Repository - The Computer Vision Foundation, accessed August 15, 2025, https://openaccess.thecvf.com/content/ICCV2023/html/Mao_COCO-O_A_Benchmark_for_Object_Detectors_under_Natural_Distribution_Shifts_ICCV_2023_paper.html
 43. MomentDiff: Generative Video Moment Retrieval from Random to Real - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2307.02869v1>
 44. Learning 2D Temporal Adjacent Networks for Moment Localization with Natural Language | Request PDF - ResearchGate, accessed August 15, 2025, https://www.researchgate.net/publication/342310309_Learning_2D_Temporal_Adjacent_Networks_for_Moment_Localization_with_Natural_Language
 45. MomentDiff: Generative Video Moment Retrieval from Random to Real - OpenReview, accessed August 15, 2025, <https://openreview.net/forum?id=SQouRKRIXY-eld=qMYsjzx1Si>
 46. [2307.02869] MomentDiff: Generative Video Moment Retrieval from Random to Real - arXiv, accessed August 15, 2025, <https://arxiv.org/abs/2307.02869>
 47. The Surprising Effectiveness of Multimodal Large Language Models for Video Moment Retrieval - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2406.18113v2>
 48. Context-Enhanced Video Moment Retrieval with Large Language Models - the UWA Profiles and Research Repository, accessed August 15, 2025, <https://research-repository.uwa.edu.au/files/368425736/2405.12540v1.pdf>
 49. Navigating the Alignment Challenges of Diffusion Models: Insights and Innovations, accessed August 15, 2025, <https://www.preprints.org/manuscript/202501.1502/v1>
 50. [2503.20516] Small Object Detection: A Comprehensive Survey on Challenges, Techniques and Real-World Applications - arXiv, accessed August 15, 2025, <https://arxiv.org/abs/2503.20516>
 51. Small Object Detection: A Comprehensive Survey on Challenges, Techniques and Real-World Applications - ResearchGate, accessed August 15, 2025, https://www.researchgate.net/publication/390213601_Small_Object_Detection_A_Comprehensive_Survey_on_Challenges_Techniques_and_Real-World_Applications
 52. Towards Small Object Editing: A Benchmark Dataset and A Training-Free Approach - arXiv, accessed August 15, 2025, <https://arxiv.org/abs/2411.01545>
 53. Length-Aware DETR for Robust Moment Retrieval - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2412.20816v1>
 54. Few-Shot Fine-Grained Image Classification: A Comprehensive Review - MDPI,

- accessed August 15, 2025, <https://www.mdpi.com/2673-2688/5/1/20>
55. A survey on generative adversarial networks for imbalance problems in computer vision tasks - PMC, accessed August 15, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC7845583/>
 56. DETR: End-to-End Object Detection With Transformers, accessed August 15, 2025, https://alcinos.github.io/detr_page/
 57. Advancing Vision-Language Models with Generative AI - Preprints.org, accessed August 15, 2025, https://www.preprints.org/frontend/manuscript/10b5ed95bd23954c58eef830d9d74bfa/download_pub
 58. (PDF) Advancing Vision-Language Models with Generative AI - ResearchGate, accessed August 15, 2025, https://www.researchgate.net/publication/388631450_Advancing_Vision-Language_Models_with_Generative_AI
 59. Vision-Language Models for Vision Tasks: A Survey - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2304.00685v2>
 60. (PDF) Transformer-Based Generative Models for Context-Aware Intrusion Detection in Zero-Day Attack Scenarios - ResearchGate, accessed August 15, 2025, https://www.researchgate.net/publication/391666745_Transformer-Based_Generative_Models_for_Context-Aware_Intrusion_Detection_in_Zero-Day_Attack_Scenarios
 61. Yi Jiang - Homepage, accessed August 15, 2025, <https://enjoyyi.github.io/>
 62. [1510.05613] PERCH: Perception via Search for Multi-Object Recognition and Localization, accessed August 15, 2025, <https://arxiv.org/abs/1510.05613>
 63. What is a Generative Model? | IBM, accessed August 15, 2025, <https://www.ibm.com/think/topics/generative-model>
 64. A Review of the State of the Art 3D Generative Models and Their Applications, accessed August 15, 2025, <https://www.ewadirect.com/proceedings/ace/article/view/17919>
 65. Recent Advances in 3D Object Detection for Self-Driving Vehicles: A Survey - MDPI, accessed August 15, 2025, <https://www.mdpi.com/2673-2688/5/3/61>
 66. A Systematic Survey of Transformer-Based 3D Object Detection for Autonomous Driving: Methods, Challenges and Trends - MDPI, accessed August 15, 2025, <https://www.mdpi.com/2504-446X/8/8/412>
 67. What the focus on multimodal, generative, and 3D AI means for the future at CVPR 2024, accessed August 15, 2025, <https://labelbox.com/blog/what-the-focus-on-multimodal-generative-and-3d-ai-mean-for-the-future-at-cvpr-2024/>
 68. A Lightweight Moment Retrieval System with Global Re-Ranking and Robust Adaptive Bidirectional Temporal Search - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2504.09298v1>
 69. ECCV 2024 Orals, accessed August 15, 2025, <https://eccv.ecva.net/virtual/2024/events/oral>
 70. Generative Models for Computer Vision - GitHub Pages, accessed August 15,

- 2025, <https://generative-vision.github.io/workshop-CVPR-25/>
71. Generative Models for Computer Vision - GitHub Pages, accessed August 15, 2025, <https://generative-vision.github.io/workshop-CVPR-24/>
72. Computer Vision Meets Generative Models in Agriculture: Technological Advances, Challenges and Opportunities - MDPI, accessed August 15, 2025, <https://www.mdpi.com/2076-3417/15/14/7663>
73. Representation Learning for Fine-Grained Change Detection - MDPI, accessed August 15, 2025, <https://www.mdpi.com/1424-8220/21/13/4486>
74. [1910.07169] Generative Modeling for Small-Data Object Detection - arXiv, accessed August 15, 2025, <https://arxiv.org/abs/1910.07169>
75. ODGEN: Domain-specific Object Detection Data Generation with Diffusion Models - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2405.15199v1>