

Training Objectives in Visual Localization: A Comprehensive Analysis of Object Detection and Video Moment Retrieval

Part I: Foundational Principles of Training Objectives in Object Detection

Object Detection (OD) stands as a cornerstone task in computer vision, requiring models to perform two fundamental operations simultaneously: identifying the class of objects within an image and precisely localizing their spatial extent. This dual requirement inherently frames object detection as a multi-task learning problem. The training objectives, therefore, are not monolithic but are composite structures designed to guide the model's optimization process across both classification and localization. The evolution of these objectives reflects a sophisticated progression from simple, computationally convenient heuristics to highly specialized functions that are closely aligned with the final evaluation metrics, addressing critical challenges such as class imbalance and the nuances of geometric accuracy.

Model Outputs as a Set of Predictions

The fundamental output of a modern object detection model is not a single prediction but a variable-length set of structured data points, where each point represents a hypothesis for a single object in the image.¹ This set-based output is a direct consequence of the task's nature, as an image can contain an arbitrary number of objects.

Deconstructing the Output Vector

Each individual prediction within the output set is typically a vector that encapsulates all necessary information about a detected object. This vector is composed of several key components:

- **Bounding Box Coordinates:** These define the spatial location of the detected object. The representation of these coordinates is a critical design choice. Two formats are prevalent:
 1. **xyxy format:** This specifies the coordinates of the top-left ($x_{\min}$, $y_{\min}$) and bottom-right ($x_{\max}$, $y_{\max}$) corners of the bounding box. This format is intuitive and widely used in frameworks and for visualization.¹
 2. **xywh format:** This specifies the coordinates of the box's center (x_{center} , y_{center}) along with its width and height. This format can be advantageous for certain regression loss calculations, as it decouples position from size.¹

These coordinates can be represented as absolute pixel values within the image's coordinate system or, more commonly, as normalized values between 0 and 1 relative to the image's dimensions. Normalization provides invariance to image size, which is crucial for robust model training.¹

- **Class Probabilities:** For each detection, the model outputs a vector of scores, where each element corresponds to one of the predefined classes in the training dataset (e.g., the 80 object categories in the COCO dataset). After applying a softmax function, this vector represents a probability distribution over the classes, indicating the likelihood that the detected object belongs to each class.¹
- **Confidence Score:** Also known as an "objectness" score, this is a scalar value between 0 and 1 that represents the model's confidence that the predicted bounding box contains *any* object, as opposed to being part of the background. During inference, this score is often multiplied by the maximum class probability to produce the final confidence score for a given detection, which is then used for ranking and thresholding.

The Role of Annotation Formats

The structure and content of the model's output are directly supervised by the ground-truth data, making the annotation format a critical component of the training pipeline. Various formats have been established, each with its own trade-offs in terms

of complexity, readability, and efficiency. Prominent examples include:

- **COCO JSON:** A comprehensive JSON-based format that has become a benchmark standard. It supports object detection, instance segmentation, and keypoint detection, encoding bounding boxes, class labels, and other metadata in a structured manner.²
- **Pascal VOC XML:** An older but still common XML-based format that is human-readable and stores bounding box coordinates and class labels for each object.³
- **YOLO TXT:** A simple and efficient text-based format where each object is represented by a single line in a text file. Each line contains the class ID and the normalized xywh coordinates of the bounding box. This minimalist format is well-suited for the efficient data loading required by YOLO-family models.¹

The diversity of these formats necessitates the use of data conversion tools to ensure compatibility between datasets and model training pipelines, highlighting a practical challenge in the field.¹

The Composite Loss: A Multi-Task Optimization Framework

To train a model to produce the complex outputs described above, the training objective is formulated as a composite loss function. This function is typically a weighted sum of two primary components: a loss for the classification task and a loss for the localization (bounding box regression) task.⁴ This structure directly mirrors the dual objectives of the detection problem.

The general mathematical formulation of this composite loss is:

$$L_{total} = \lambda_{cls} L_{cls} + \lambda_{loc} L_{loc}$$

Here, L_{cls} quantifies the error in the predicted class probabilities, while L_{loc} measures the discrepancy between the predicted and ground-truth bounding box coordinates. The hyperparameters λ_{cls} and λ_{loc} are crucial for balancing the influence of the two tasks during training. An improper balance can lead to a model that excels at classification but fails to localize objects accurately, or vice versa.⁷

The loss function serves as the mathematical heart of the learning process. It provides a quantitative measure of the model's performance on the training data. During training, the gradient of this loss function with respect to the model's

parameters (weights and biases) is calculated. This gradient is then used by an optimization algorithm, such as stochastic gradient descent, to update the parameters via backpropagation. This iterative process of calculating the loss and updating the weights aims to find a set of parameters that minimizes the total loss, thereby improving the model's predictive accuracy.⁴

Classification Loss: Addressing Class Imbalance

The classification component of the loss function is responsible for training the model to correctly identify the category of each detected object. While standard classification losses form the foundation, their direct application in object detection reveals a significant challenge unique to the task.

The Standard: Cross-Entropy Loss

The conventional choice for multi-class classification is the Cross-Entropy (CE) loss. It measures the dissimilarity between the predicted probability distribution over classes and the true distribution, which is typically a one-hot vector representing the ground-truth class.⁶ For a single prediction, the CE loss is calculated as:

$$L_{CE}(p, y) = -\sum_{c=1}^C y_c \log(p_c)$$

where C is the number of classes, y_c is 1 if c is the true class and 0 otherwise, and p_c is the model's predicted probability for class c .

The Challenge of Class Imbalance

A defining characteristic of object detection, particularly in single-stage detectors that evaluate a dense grid of potential object locations (anchors), is the extreme imbalance between foreground (objects) and background classes. For any given image, the vast majority of anchors will not contain an object and will be labeled as background. This creates a scenario where "easy negative" examples (correctly

classified background regions) overwhelm the training process.¹¹ When using standard Cross-Entropy loss, the accumulated loss from these thousands of easy negatives can dominate the total loss, dwarfing the contribution from the few, more informative "hard positive" and "hard negative" examples. This causes the model's optimization to focus on correctly classifying the ubiquitous background, leading to poor learning of features necessary to detect the rare foreground objects.¹²

The Evolution: Focal Loss

The invention of Focal Loss was a direct and elegant solution to this debilitating class imbalance problem. It was a key innovation that enabled single-stage detectors like RetinaNet to surpass the accuracy of more complex two-stage detectors for the first time.¹¹ Focal Loss modifies the standard Cross-Entropy loss by introducing a modulating factor

$(1-p_t)^\gamma$, where p_t is the model's estimated probability for the ground-truth class and γ is a tunable focusing parameter.

The mathematical formulation of Focal Loss is:

$$\text{LFL}(p_t) = -\alpha t (1-p_t)^\gamma \log(p_t)$$

The modulating factor has a powerful effect:

- For well-classified examples where p_t is high (e.g., an easy background anchor with $p_t \rightarrow 1$), the $(1-p_t)^\gamma$ term approaches zero, effectively down-weighting the loss contribution from this example.
- For poorly-classified examples where p_t is low (e.g., a hard-to-detect object), the $(1-p_t)^\gamma$ term approaches one, preserving its contribution to the loss.

By dynamically scaling the loss based on classification confidence, Focal Loss forces the training process to focus on the small set of hard examples, mitigating the dominance of easy negatives without requiring explicit hard negative mining or complex sampling strategies.¹² This demonstrates a powerful principle in model development: sometimes, the most significant performance gains come not from architectural changes, but from a fundamental redesign of the objective function to address the specific statistical challenges of the task.

Bounding Box Regression Loss: Measuring Localization Accuracy

The localization component of the loss function trains the model to predict bounding box coordinates that align as closely as possible with the ground-truth boxes. The evolution of this loss component showcases a clear trend towards functions that are more directly aligned with the final evaluation metric for localization, Intersection over Union (IoU).

The L-Norm Family (Distance-based)

Initial object detection models relied on standard regression losses from the L-norm family to penalize differences between the predicted and true bounding box coordinate vectors.

- **L2 Loss (Mean Squared Error):** Calculates the average of the squared differences between coordinates. It is sensitive to large errors (outliers) due to the squaring term.⁸
- **L1 Loss (Mean Absolute Error):** Calculates the average of the absolute differences. It is more robust to outliers than L2 loss but can have unstable gradients.¹⁶
- **Smooth L1 Loss (Huber Loss):** This became the de facto standard for many years. It behaves like L2 loss for small errors (providing stable gradients near the optimum) and like L1 loss for large errors (providing robustness to outliers). Its formulation is a piecewise function ⁶:

$$L_{smooth\ L1}(x) = \begin{cases} 0.5x^2 & |x| < 1 \\ |x| - 0.5 & \text{otherwise} \end{cases}$$

A significant drawback of all L-norm losses is their underlying assumption that the four coordinate parameters (x,y,w,h) are independent. Minimizing the distance for each parameter separately does not necessarily maximize the geometric overlap of the boxes, which is what the IoU metric actually measures. This created a misalignment between the training objective and the evaluation metric.¹⁷

The IoU Family (Metric-aligned)

To address the misalignment of L-norm losses, a new family of loss functions based directly on the Intersection over Union (IoU) metric was developed. This shift represents a move towards the "metric-alignment principle," where the training objective is designed to more closely reflect the ultimate performance measure.

- **IoU Loss:** The most direct approach is to use the IoU value itself as a loss, typically formulated as $L_{IoU}=1-IoU$. This directly optimizes for the desired metric. However, it suffers from two major flaws: (1) if the predicted and ground-truth boxes do not overlap, IoU is 0, and the loss becomes 1, resulting in a zero gradient. This means the model receives no learning signal to move the boxes closer together. (2) It cannot distinguish between different degrees of non-overlap; a box that is far away and one that is just touching both have an IoU of 0.¹⁰
- **Generalized IoU (GIoU) Loss:** GIoU was the first major improvement, designed to solve the non-overlapping problem. It introduces a penalty term based on the area of the smallest convex box that encloses both the predicted and ground-truth boxes (C). The loss is formulated as $LGIoU=1-IoU+|C|-|C-(A \cup B)|$, where A and B are the predicted and ground-truth boxes. This penalty term ensures that there is always a non-zero gradient, pushing non-overlapping boxes towards each other. While effective, GIoU can converge slowly and may result in inaccurate regressions, as it sometimes prioritizes expanding the predicted box to cover the ground truth over achieving a tight fit.¹⁷
- **Distance IoU (DIoU) Loss:** DIoU improves upon GIoU by directly incorporating the distance between the center points of the two boxes. It adds a penalty term that minimizes the normalized Euclidean distance between the box centers, leading to significantly faster convergence and more accurate localization.¹⁹
- **Complete IoU (CIoU) Loss:** CIoU further refines DIoU by adding an additional penalty term that accounts for the consistency of the aspect ratios between the predicted and ground-truth boxes. By considering three key geometric factors—overlap area, center point distance, and aspect ratio—CIoU provides a more comprehensive measure of box similarity, resulting in faster convergence and higher-quality bounding box predictions.¹⁰

The following table provides a comparative summary of these influential bounding box regression losses.

Loss Function	Core Idea	Advantages	Disadvantages
---------------	-----------	------------	---------------

Smooth L1 Loss	A hybrid of L1 and L2 loss applied to coordinate differences.	Robust to outliers; stable gradients near the optimum.	Treats coordinates as independent; not directly aligned with the IoU metric. ¹⁷
GIoU Loss	Penalizes non-overlap based on the area of the smallest enclosing box.	Provides a gradient for non-overlapping boxes. ¹⁹	Slow convergence; can lead to inaccurate box shapes by prioritizing area reduction over overlap. ¹⁹
DIoU Loss	Adds a penalty for the normalized distance between box centers.	Faster convergence than GIoU; directly minimizes center point distance. ¹⁹	Does not consider the aspect ratio of the boxes.
CIoU Loss	Adds a penalty for aspect ratio consistency to DIoU.	Fastest convergence; considers overlap, distance, and shape, leading to more accurate regressions. ¹⁹	Can be slightly more complex to implement than simpler IoU variants.

Advanced Objectives: Ranking-Based Loss Functions

The most recent advancements in loss function design take the metric-alignment principle to its logical conclusion. Instead of optimizing proxies for the final evaluation metric, Average Precision (AP), these methods attempt to create differentiable surrogates for the AP metric itself. This is challenging because AP, which depends on the ranked order of predictions, is inherently non-differentiable.

- AP Loss and aLRP Loss:** These pioneering works reframe the classification task as a ranking problem. The objective is to ensure that the scores of positive predictions are ranked higher than the scores of all negative predictions. They address the non-differentiable components of the AP calculation, such as the ranking operation, by using error-driven update rules inspired by perceptron learning.⁷ This approach more directly optimizes the quantity of interest and has been shown to be more robust to the severe class imbalance inherent in object detection.²¹

- **Rank & Sort (RS) Loss:** This method extends the ranking concept further. The RS Loss objective has two components: (1) **Rank:** Every positive prediction must be ranked above every negative prediction. (2) **Sort:** The positive predictions should be sorted among themselves according to their localization quality (IoU). A positive prediction with a higher IoU should receive a higher confidence score from the model. This sorting objective explicitly encourages the model to be more confident about its better-localized predictions, which directly improves the AP score. A key advantage of these ranking-based losses is that they often require fewer hyperparameters and can handle extreme class imbalance without specialized sampling heuristics.⁷

Part II: The Transformer Revolution: DETection TRansformer (DETR)

The introduction of the DETection TRansformer (DETR) by Facebook AI Research in 2020 marked a significant paradigm shift in object detection. DETR reconceptualized the task not as a dense classification and regression problem, but as a direct set prediction problem. This approach eliminated the need for many of the complex, hand-crafted components that had defined object detectors for years, such as anchor boxes and Non-Maximum Suppression (NMS), offering a truly end-to-end solution.²³

A New Paradigm: Object Detection as Direct Set Prediction

The core philosophy of DETR is to treat object detection as a set-to-set translation task: given an image, the model directly outputs a finite set of object predictions. This elegant formulation removes the intermediate steps of generating thousands of proposals or anchors and then filtering them, which were central to previous architectures.²⁵

Architecture Overview

The DETR architecture combines a standard convolutional backbone with a Transformer encoder-decoder, a structure originally developed for natural language processing tasks.²³

1. **CNN Backbone:** A convolutional neural network (e.g., ResNet) is used to extract a rich feature map from the input image. This feature map is then flattened and supplemented with positional encodings to provide the Transformer with spatial information, which it otherwise lacks.²³
2. **Transformer Encoder:** The flattened feature map is passed through a Transformer encoder. Using its multi-head self-attention mechanism, the encoder models the global context of the image, capturing complex relationships between all parts of the visual scene. This allows it to reason about the entire image content at once.²³
3. **Transformer Decoder and Object Queries:** The decoder is where DETR's novelty is most apparent. Instead of processing the encoder's output autoregressively (one token at a time), it takes two inputs: the processed image features from the encoder and a small, fixed-size set of learned embeddings called **object queries**.²³ These object queries (typically $N=100$) are learned during training and act as slots or "hooks" for potential objects. The decoder uses self-attention and cross-attention mechanisms to transform these queries into the final object predictions. Each output embedding from the decoder is then passed through a shared feed-forward network (FFN) to predict either a class label and a bounding box, or a "no object" class.²⁶ This parallel decoding process is a key architectural feature that contributes to DETR's efficiency.²⁶

The Bipartite Matching Loss: A Unique Assignment Problem

The set-based prediction approach introduces a unique challenge for the loss function. The model outputs an unordered set of N predictions, while the ground truth consists of a smaller, unordered set of M objects (where $M < N$). To compute a loss, a one-to-one correspondence must be established between a subset of the N predictions and the M ground-truth objects. Predictions that are not matched to any ground-truth object should be penalized if they do not predict the "no object" class.²⁹

The Hungarian Algorithm

DETR solves this assignment problem using the **Hungarian algorithm**, a classic combinatorial optimization algorithm that can find the minimum weight matching in a bipartite graph.²⁹ In this context, the two sets of nodes in the graph are the model's

N predictions and the ground truth's M objects (padded with $N-M$ "no object" slots). The weight of an edge between a prediction and a ground-truth object is a **pairwise matching cost**. This cost is defined to encourage good matches and is typically a linear combination of:

- The negative log-probability of the correct class.
- A bounding box similarity loss (e.g., a combination of L1 loss and GloU loss).²⁹

The Hungarian algorithm efficiently finds the permutation of predictions, σ^* , that minimizes the total matching cost across all ground-truth objects.³⁰

The Set-Based Global Loss (Hungarian Loss)

Once the optimal one-to-one matching is established by the Hungarian algorithm, the final loss, referred to as the **Hungarian Loss**, can be computed. This loss is calculated as a sum over only the matched pairs. For each matched pair $(i, \sigma^*(i))$, where i is the index of a ground-truth object and $\sigma^*(i)$ is the index of the matched prediction, the loss is a sum of a classification term and a localization term.²⁸

The mathematical formulation can be summarized in two steps:

1. Find the optimal matching σ^* :

$$\sigma^* = \sigma \in S_N \arg \min_i \sum_N L_{\text{match}}(y_i, y^{\sigma(i)})$$

where S_N is the set of all permutations of N elements, y_i is the i -th ground-truth object (padded with $\emptyset$), $y^{\sigma(i)}$ is the prediction, and L_{match} is the pairwise matching cost.

2. Compute the Hungarian loss:

$$L_{\text{Hungarian}}(y, y^*) = \sum_{i=1}^N [-\log p^{\sigma^*(i)}(c_i) + 1_{\{c_i \neq \emptyset\}} L_{\text{box}}(b_i, b^{\sigma^*(i)})]$$

where c_i and b_i are the class label and box of the ground-truth object y_i , and $\hat{p}$ and $\hat{b}$ are the predicted class probability and box. The indicator function $1_{\{c_i \neq \emptyset\}}$ ensures that the box loss is only applied for actual objects, not the padded "no object" slots. The loss for the "no object" class is typically down-weighted to account for the class imbalance.³⁰

This elegant combination of a set-prediction architecture and a bipartite matching loss allowed DETR to remove the complex, heuristic-based components of its predecessors. This represented the culmination of a long-standing goal in the field: to create a truly end-to-end, learnable system for object detection. However, this architectural simplicity came at a price. The global attention mechanism was computationally intensive, and the model required significantly longer training schedules to converge compared to CNN-based detectors.²⁸ This initial inefficiency was a direct catalyst for a subsequent wave of research focused on optimizing the Transformer for detection, leading to more practical and powerful models like RT-DETR, which introduced an efficient hybrid encoder to achieve real-time performance.³² This cycle—a revolutionary but initially impractical idea followed by efficiency-focused refinements—is a recurring pattern in the advancement of deep learning.

Part III: Training Objectives in Video Moment Retrieval

Shifting from the spatial domain of images to the temporal domain of videos, Video Moment Retrieval (VMR) presents a distinct set of challenges centered on understanding dynamic content and its relationship with natural language. Also known as Temporal Sentence Grounding in Videos (TSGV) or Natural Language Video Localization (NLVL), the task requires a model to identify and localize a specific temporal segment, or "moment," within an untrimmed video that corresponds to a given linguistic query.³⁴

The Task of Temporal Grounding

The fundamental challenge in VMR is **cross-modal alignment**: bridging the semantic

gap between the continuous, high-dimensional visual data of a video and the discrete, symbolic information of a text query.³⁶ A successful model must learn a joint representation or an effective interaction mechanism that maps words and phrases like "a person slicing a cucumber" to the specific sequence of frames where that action occurs. This requires not only recognizing objects and actions but also understanding their temporal relationships and context.³⁹

Model Output Formats for Temporal Localization

The architectural approaches to VMR are broadly divided into two families, proposal-based and proposal-free, which directly dictate the format of the model's output and its training objective.⁴²

Proposal-Based Methods

These methods adopt a two-stage "propose-and-rank" pipeline, analogous to early two-stage object detectors.

- **Process:** First, a large number of candidate temporal segments (proposals) are generated. This can be done using multi-scale sliding windows across the video or through more sophisticated, learnable proposal generation networks.³⁶ In the second stage, each of these proposals is compared against the text query to compute a relevance or similarity score.
- **Output:** The final output is a ranked list of temporal moments, each defined by a start time, an end time, and a confidence score indicating its semantic alignment with the query.⁴⁶ A key advantage of this approach is its natural ability to handle multi-target VMR, where a single query may correspond to multiple distinct moments in a video, as it can return all proposals that exceed a certain score threshold.⁴⁴

Proposal-Free (Regression-Based) Methods

Aiming for greater efficiency and an end-to-end architecture, proposal-free methods eliminate the explicit proposal generation step.

- **Process:** These models directly predict the boundaries of the target moment from the fused video and query features.
- **Output:** The output can take two primary forms:
 1. **Direct Boundary Regression:** The model outputs two scalar values representing the predicted start and end timestamps of the moment.⁴³
 2. **Temporal Probability Maps:** The model predicts two probability distributions over the video's duration. One map represents the probability of each frame being the *start* of the target moment, and the other represents the probability of it being the *end*. The final moment is then determined by finding the pair of start and end frames with the highest joint probability.⁴⁴

Proposal-free methods are generally faster and computationally less expensive, but most are inherently designed to predict a single moment, making them less suitable for multi-target scenarios.⁴³

The following table compares these two dominant methodologies in VMR.

Methodology	Core Approach	Typical Model Output	Advantages	Disadvantages	Suitability
Proposal-Based	Two-stage "propose-and-rank": generate candidate moments, then score each against the query.	A ranked list of moments, each with (start, end, score). ⁴⁶	Naturally handles multi-target retrieval; can achieve high recall with dense proposals. ⁴⁴	Computationally expensive; performance depends on the quality of the initial proposals. ⁴⁸	Multi-target VMR, scenarios requiring high precision.
Proposal-Free	End-to-end prediction of moment boundaries without explicit candidates.	Direct regression of (start, end) timestamps or two temporal probability	Fast and efficient; simpler end-to-end architecture.	Typically designed to predict a single moment; may struggle with multi-target	Real-time applications, single-target VMR.

		maps. ⁴³		queries. ⁴⁴	
--	--	---------------------	--	------------------------	--

Loss Functions for Cross-Modal Alignment

The core of training a VMR model is teaching it to associate visual and textual concepts. This is primarily achieved through loss functions that encourage the model to learn a shared embedding space where semantically related video and text representations are brought closer together.

The Central Role of Contrastive Learning

Contrastive learning is the dominant paradigm for this task. The underlying principle is to train the model to distinguish between "positive" (matching) video-query pairs and "negative" (non-matching) pairs.⁴⁹

- InfoNCE (Noise-Contrastive Estimation):** A widely used formulation of contrastive loss is InfoNCE. For a given positive pair (v,q), the loss function aims to maximize their similarity score relative to the scores of a set of negative samples. The negative samples can be other queries from the same batch (in-batch negatives) or other video segments.⁴⁴ The typical mathematical formulation for a single positive pair is:

$$L_{\text{contrastive}} = -\log(\exp(\text{sim}(v,q)/\tau) + \sum_{i=1}^K \exp(\text{sim}(v,q_{\text{negi}})/\tau) \exp(\text{sim}(v,q)/\tau))$$
 Here, $\text{sim}(\cdot, \cdot)$ is a similarity function like cosine similarity, τ is a temperature hyperparameter that controls the sharpness of the distribution, and $\{q_{\text{negi}}\}$ is the set of K negative query samples. A similar term can be constructed for negative video samples. Other variants, such as the triplet margin loss or hinge loss, are also used to enforce a margin between the scores of positive and negative pairs.⁴⁷
- The "Weak Alignment" Problem:** A significant challenge in VMR is that the ground-truth annotations are often imprecise or ambiguous. A query might describe a complex event, but the annotated start and end times may be subjective or include irrelevant frames at the boundaries.⁴² This "weak alignment" provides a noisy supervision signal, making it difficult for models to learn precise semantic correspondences. This data-inherent problem has driven the development of more robust training strategies, such as weakly-supervised

methods that learn from video-level labels instead of expensive boundary annotations, or fine-grained contrastive losses that operate at the frame-word level to identify the most salient content within a moment.³⁵

Loss Functions for Temporal Boundary Regression

In addition to aligning semantics, the model must also be trained to predict accurate temporal boundaries. This regression task is directly analogous to bounding box regression in object detection, and the loss functions are adapted accordingly. The dual nature of VMR—being both a retrieval (alignment) and a localization (regression) task—is reflected in the use of composite loss functions that combine objectives for both aspects.

- **Adapting Spatial Losses for the Temporal Domain:**
 - **L1/L2 Loss:** Many VMR models employ a regression loss, such as L1 or Smooth L1 loss, on the predicted temporal boundaries. The loss is calculated based on the difference between the predicted (start, end) timestamps and the ground-truth (start_gt, end_gt) timestamps.⁴⁵
 - **Temporal IoU (tIoU) Loss:** Following the metric-alignment principle, a more effective approach is to use a loss based on the temporal Intersection over Union (tIoU). The tIoU is the ratio of the length of the intersection of the predicted and ground-truth segments to the length of their union. The loss is typically formulated as $L_{tIoU} = 1 - tIoU$. This directly optimizes the primary metric used for evaluating localization performance in VMR.⁵³
- **Hybrid Objectives:** A common training strategy is to use a weighted sum of an alignment loss and a regression loss. For example, a proposal-based method might be trained with a total loss function like:
$$L_{total} = \lambda_{align} L_{align} + \lambda_{reg} L_{reg}$$

Here, L_{align} could be a binary cross-entropy or contrastive loss on the proposal scores, while L_{reg} could be a tIoU or Smooth L1 loss on the coordinates of the highest-scoring positive proposals.⁴⁷ This hybrid objective ensures that the model learns both *what* to look for (semantic alignment) and *where* to find it precisely (temporal localization).

Part IV: Adapting Transformers for Temporal Localization

The success of the DETR architecture in revolutionizing 2D object detection naturally led to its adaptation for 1D temporal localization tasks like Video Moment Retrieval and Temporal Action Detection (TAD). This required re-engineering the model's inputs, outputs, and positional encodings for the temporal domain, while preserving the core principles of end-to-end set prediction and bipartite matching.

Moment-DETR: An End-to-End Framework for VMR

Moment-DETR was a pioneering work that applied the DETR philosophy directly to VMR, creating an end-to-end framework for jointly performing moment retrieval and highlight detection (a related task of predicting a saliency score for every clip).⁵⁴

Architectural Adaptations

The adaptation from a 2D spatial model to a 1D temporal model involved several key modifications:

- **Input Representation:** Instead of processing a 2D feature map from a CNN, Moment-DETR takes a sequence of 1D feature vectors as input, where each vector represents a short clip or segment of the video.⁵⁴
- **Encoder-Decoder Structure:** The Transformer encoder-decoder architecture is largely retained. The encoder uses self-attention to model long-range temporal dependencies across the entire video. The decoder takes the encoder's output and a set of learnable "moment queries" to predict a set of potential moments.⁵⁶
- **Prediction Head:** The feed-forward network (FFN) at the output of the decoder is modified. Instead of predicting a 2D bounding box (x, y, w, h) and a class label, it predicts a 1D temporal segment, typically parameterized by its center and duration (center, duration), along with a saliency score.⁵⁸
- **Loss Function:** The training objective remains a set-based bipartite matching loss. The Hungarian algorithm is used to find the optimal assignment between predicted moments and ground-truth moments. The final loss is a combination of

a temporal localization loss (a mix of L1 loss and tIoU loss) and a saliency loss (e.g., Binary Cross-Entropy) for the highlight detection task.⁶⁰

This successful adaptation demonstrates the power and generality of the DETR framework. The core concepts of abstracting the problem as a set prediction task and using a bipartite matching loss are not inherently tied to 2D spatial data. By re-engineering the low-level components—such as the input/output representations and positional encodings—the high-level conceptual framework can be effectively transferred to a new domain with different dimensionality.

Architectural Evolution: Insights from DualDETR for Temporal Action Detection

While VMR and Temporal Action Detection (TAD) are distinct tasks—VMR localizes moments based on free-form text, while TAD localizes and classifies actions from a predefined set—innovations in one area often provide valuable insights for the other. The development of DualDETR for TAD highlights a key limitation of directly applying the standard DETR paradigm to temporal tasks and offers a sophisticated solution.⁶¹

The Problem with Instance-Level Prediction

A direct adaptation of DETR, like Moment-DETR, typically predicts a temporal segment at the "instance level" by regressing its center and duration. While this captures the holistic content of the moment, it can lead to sub-optimal boundary localization. The reason for this is a fundamental difference between objects and actions: an object's location is well-defined by its center, but a temporal action or event is often better defined by its start and end points, which may correspond to distinct visual cues (e.g., an object being picked up vs. being put down).⁶¹ Asking a single set of decoder queries to simultaneously capture the overall semantic content and the precise boundary transitions is a difficult task.

DualDETR's Solution: Two-Branch Decoding

DualDETR addresses this challenge by proposing a novel two-branch decoding structure that decouples the localization task into two levels of granularity ⁶¹:

1. **Instance-Level Branch:** One set of decoder queries operates as in a standard DETR, predicting the action instance as a whole (center and duration). This branch excels at capturing the main semantic content of the action.
2. **Boundary-Level Branch:** A second, separate set of queries is dedicated to predicting the start and end boundaries directly. These queries are trained to be more sensitive to the fine-grained temporal details and transitions that mark the edges of an action.

These two branches use distinct decoders, allowing them to learn specialized representations for their respective sub-tasks. The predictions from both branches are aligned through a joint query initialization strategy, enabling them to work in a complementary fashion to produce a final prediction that is both semantically accurate and temporally precise. ⁶¹

The evolution from a single-branch architecture in Moment-DETR to a more specialized two-branch architecture in DualDETR illustrates an important principle: as the demand for precision in a complex task increases, monolithic, general-purpose architectures often give way to more modular designs. By decomposing the problem into sub-tasks with different semantic requirements (content identification vs. boundary localization) and assigning specialized "expert" components to each, models can achieve a higher level of performance. This suggests a future trend in temporal localization where complex reasoning will be handled by a coordinated system of specialized modules within a unified framework.

Part V: Synthesis and Future Directions

The fields of Object Detection and Video Moment Retrieval, while addressing localization in different domains (spatial vs. temporal), have followed remarkably parallel evolutionary paths. Both have progressed from complex, multi-stage pipelines with heuristic-based objectives to streamlined, end-to-end architectures trained with metric-aligned loss functions. The introduction of the Transformer has served as a powerful unifying force, providing a common architectural language for solving these

set prediction problems. Looking forward, the trajectory of both fields points towards a future dominated by large-scale foundation models, with a growing emphasis on robustness, efficiency, and the ability to handle the complexity of real-world scenarios.

A Comparative Analysis: Spatial vs. Temporal Localization

While both tasks fall under the umbrella of "visual localization," their fundamental differences in domain and objective lead to distinct challenges and solutions.

- **Fundamental Similarities:**
 - **Multi-Task Nature:** Both tasks require a combination of classification/alignment and localization/regression. This is reflected in their composite loss functions.
 - **Metric-Alignment:** Both have seen performance gains from shifting their regression losses from simple L-norm distances to functions based on Intersection over Union (IoU), which better reflects the evaluation metric.¹⁷
 - **End-to-End Learning:** The Transformer architecture has enabled both fields to move towards fully end-to-end solutions (DETR, Moment-DETR), eliminating complex post-processing steps.²⁵
- **Fundamental Differences:**
 - **Dimensionality:** OD is a 2D spatial problem concerned with (x, y) position and scale within a static image. VMR is a 1D temporal problem concerned with (start, end) position and duration within a dynamic video sequence.⁶⁶
 - **Core Challenge:** In OD, a primary challenge is managing the extreme class imbalance and precisely localizing a potentially large number of objects from a fixed set of classes. In VMR, the central challenge is bridging the **cross-modal semantic gap** between an open-vocabulary, free-form language query and the visual content of a video.³⁷
 - **Nature of Context:** Spatial context in OD involves the relationships between objects within a scene (e.g., a person is *on* a horse). Temporal context is paramount in VMR, as the meaning of a moment is often defined by the sequence of events that precede and follow it.⁶⁸
 - **Feature Extraction:** OD models typically use 2D CNNs to extract spatial features. VMR models must capture motion and temporal dynamics, often employing 3D CNNs or a combination of 2D CNNs with recurrent or temporal attention mechanisms.⁷⁰

Emerging Trends and Unsolved Challenges (2024-2025)

The current research frontier in both fields is moving beyond standard benchmark performance to address the complexities of real-world application.

- **The Impact of Large Language Models (LLMs):** The most significant recent trend is the integration of LLMs and, more broadly, Vision-Language Models (VLMs). In VMR, LLMs are being used to generate rich, descriptive textual narratives for videos. These narratives serve as a bridge, enhancing the video representation with high-level semantic context and enabling the model to better understand and localize complex, nuanced queries that involve reasoning about environments, characters, and actions.⁷¹ This points towards a future where localization is not just a pattern matching task but a deep, context-aware comprehension task.
- **Data-Centric AI and Open-Vocabulary Detection:** In object detection, there is a growing focus on data-centric approaches. This includes the development of methods to train detectors on synthetic data generated by diffusion models (e.g., InstaGen), which can help expand the number of detectable categories and improve performance in data-sparse scenarios.⁷⁴ Concurrently, models like YOLO-WORLD are pushing towards real-time, open-vocabulary detection by leveraging large-scale vision-language pre-training, allowing them to detect objects from categories not seen during training.³³ This marks a shift from closed-set to open-set recognition, a crucial step towards more general-purpose vision systems.
- **Long-Form Video and Efficiency:** A major unsolved challenge for VMR is scaling to long-form videos, such as hour-long tutorials or movies. Most current benchmarks consist of videos that are only a few minutes long.⁷⁵ Processing long videos introduces immense computational and memory challenges. Emerging training strategies are exploring coarse-to-fine approaches, where an efficient model first retrieves relevant keyframes or coarse segments, which are then processed by a more powerful model for precise localization.⁷⁷
- **Robustness, Faithfulness, and Ambiguity:** As models become more accurate on clean benchmark data, the research focus is shifting to their reliability in the wild. This has led to the development of new tasks and training objectives designed to handle real-world complexities:
 - **Negative Queries:** A robust VMR system should not hallucinate a result when

a query describes an event that is not present in the video. Research on "open-set VMR" and negative query rejection is developing models that can gracefully say "I don't know".⁴⁹

- **Ambiguous Queries:** Real-world queries are often ambiguous. Novel loss functions and training strategies are being designed to handle this, for example, by calibrating contrastive losses to account for "soft positive" samples (moments that are partially relevant but not perfect matches) or by using interactive frameworks to resolve ambiguity.⁵¹
- **Faithfulness Evaluation:** New benchmarks and evaluation frameworks (e.g., MVMR) are being proposed to test a model's faithfulness by presenting it with a massive set of "distractor" videos, forcing it to not only find a moment but also to be confident that it is in the correct video.⁸²

Concluding Remarks

The journey of training objectives in object detection and video moment retrieval has been one of increasing sophistication and alignment with real-world goals. The initial reliance on simple, differentiable proxies for error has given way to complex, metric-aware loss functions that directly optimize for performance. The architectural paradigm has converged on the Transformer, whose principles of attention and set prediction have proven remarkably generalizable across both spatial and temporal domains.

The future of these fields appears to be converging towards a unified foundation model paradigm. Instead of training specialized models from scratch, future work will likely focus on adapting large, pre-trained Vision-Language Models for specific localization tasks. The training objectives will evolve accordingly, shifting from full-scale training to efficient fine-tuning, prompt engineering, and the design of lightweight adapter modules. The ultimate goal is no longer just accuracy on a closed-set benchmark, but the creation of robust, faithful, and generalizable systems that can understand and localize content in the vast and unstructured world of visual data, responding intelligently to the full spectrum of human queries. The distinction between spatial and temporal localization will likely continue to blur, subsumed by holistic video comprehension frameworks that can reason seamlessly about what is happening, where it is, and when it occurs.

Works cited

1. Object Detection - Ultralytics YOLO Docs, accessed August 15, 2025, <https://docs.ultralytics.com/tasks/detect/>
2. Object detection - Hugging Face, accessed August 15, 2025, https://huggingface.co/docs/transformers/tasks/object_detection
3. Computer Vision Annotation Formats - Roboflow, accessed August 15, 2025, <https://roboflow.com/formats>
4. Loss Functions in Deep Learning: A Comprehensive Review - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2504.04242v1>
5. A Visual Guide to Loss Functions in Object Detection and Segmentation - Medium, accessed August 15, 2025, <https://medium.com/@avithaljunk/a-visual-guide-to-loss-functions-in-object-detection-and-segmentation-fe04f75082a6>
6. LOSS FUNCTION DISCOVERY FOR OBJECT DETECTION VIA CONVERGENCE-SIMULATION DRIVEN SEARCH - OpenReview, accessed August 15, 2025, <https://openreview.net/pdf?id=5jzlpHvvRk>
7. Rank & Sort Loss for Object Detection and Instance Segmentation - CVF Open Access, accessed August 15, 2025, https://openaccess.thecvf.com/content/ICCV2021/papers/Oksuz_Rank_Sort_Loss_for_Object_Detection_and_Instance_Segmentation_ICCV_2021_paper.pdf
8. What is Loss Function? | IBM, accessed August 15, 2025, <https://www.ibm.com/think/topics/loss-function>
9. Loss Functions in Deep Learning: A Comprehensive Review - arXiv, accessed August 15, 2025, <https://arxiv.org/pdf/2504.04242>
10. Loss Function in AI/ML Explained - Ultralytics, accessed August 15, 2025, <https://www.ultralytics.com/glossary/loss-function>
11. Focal Loss for Dense Object Detection - PubMed, accessed August 15, 2025, <https://pubmed.ncbi.nlm.nih.gov/30040631/>
12. Cross Entropy loss and its Evolution to Focal loss in Object ... - Medium, accessed August 15, 2025, <https://medium.com/@abhishekjainindore24/cross-entropy-loss-and-its-evolution-to-focal-loss-in-object-detection-3dd5e28c571e>
13. A Method for Constructing a Loss Function for Multi-Scale Object Detection Networks - PMC, accessed August 15, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11946842/>
14. rmoklesur.medium.com, accessed August 15, 2025, <https://rmoklesur.medium.com/reasons-to-choose-focal-loss-over-cross-entropy-5fdccb25d282#:~:text=Reduces%20the%20impact%20of%20easy.more%20importance%20to%20difficult%20examples.>
15. What is Focal Loss and when should you use it? - Aman Arora's Blog, accessed August 15, 2025, <https://amaarora.github.io/posts/2020-06-29-FocalLoss.html>
16. Loss Functions in Machine Learning Explained - DataCamp, accessed August 15, 2025, <https://www.datacamp.com/tutorial/loss-function-in-machine-learning>
17. Smooth GloU Loss for Oriented Object Detection in Remote Sensing Images -

- MDPI, accessed August 15, 2025, <https://www.mdpi.com/2072-4292/15/5/1259>
18. Smooth GloU Loss for Oriented Object Detection in Remote Sensing Images, accessed August 15, 2025, https://www.researchgate.net/publication/368868796_Smooth_GIoU_Loss_for_Oriented_Object_Detection_in_Remote_Sensing_Images
 19. GloU, CloU and DIoU: Variants of IoU and how they are better ..., accessed August 15, 2025, <https://medium.com/@abhishekjainindore24/giou-ciou-and-diou-variants-of-iou-and-how-they-are-better-compared-to-iou-4610a015643a>
 20. Generalized IoU loss for Object Detection with Torchvision | Towards Data Science, accessed August 15, 2025, <https://towardsdatascience.com/generalized-iou-loss-for-object-detection-with-torchvision-9534029d1a89/>
 21. Searching Parameterized AP Loss for Object Detection, accessed August 15, 2025, <https://proceedings.neurips.cc/paper/2021/file/b9009beb804fa097c04d226a8ba5102e-Paper.pdf>
 22. Bucketed Ranking-based Losses for Efficient Training of Object Detectors - European Computer Vision Association, accessed August 15, 2025, https://www.ecva.net/papers/eccv_2024/papers_ECCV/papers/07634.pdf
 23. What is DETR (Detection Transformers)? - Roboflow Blog, accessed August 15, 2025, <https://blog.roboflow.com/what-is-detr/>
 24. facebookresearch/detr: End-to-End Object Detection with Transformers - GitHub, accessed August 15, 2025, <https://github.com/facebookresearch/detr>
 25. Introduction to DETR (Detection Transformers): Everything You Need to Know - Lightly, accessed August 15, 2025, <https://www.lightly.ai/blog/detr>
 26. End-to-End Object Detection with Transformers, accessed August 15, 2025, <http://arxiv.org/pdf/2005.12872>
 27. DETR - Hugging Face, accessed August 15, 2025, https://huggingface.co/docs/transformers/model_doc/detr
 28. DETR:End-to-End Object Detection with Transformers | Learning-Deep-Learning, accessed August 15, 2025, https://patrick-lgc.github.io/Learning-Deep-Learning/paper_notes/detr.html
 29. Set Prediction Loss for Object Detection | by Ненавин | Jul, 2025 - Medium, accessed August 15, 2025, <https://medium.com/@natsunoyuki/set-prediction-loss-for-object-detection-6af1f06982bb>
 30. Papers Explained 79: DETR. DETection TRansformer or DETR... | by Ritvik Rastogi, accessed August 15, 2025, <https://ritvik19.medium.com/papers-explained-79-detr-bcdd53355d9f>
 31. Introduction to DETR - Part 2: The Crucial Role of the Hungarian Algorithm | DigitalOcean, accessed August 15, 2025, <https://www.digitalocean.com/community/tutorials/introduction-detr-hungarian-algorithm-2>
 32. Baidu's RT-DETR: A Vision Transformer-Based Real-Time Object Detector,

- accessed August 15, 2025, <https://docs.ultralytics.com/models/rtdetr/>
33. CVPR 2024: Top Artificial Intelligence and Computer Vision Papers Accepted - Encord, accessed August 15, 2025, <https://encord.com/blog/cvpr-2024/>
 34. www3.ntu.edu.sg, accessed August 15, 2025, <https://www3.ntu.edu.sg/home/axsun/slides/VideoMomentRetrieval.pdf>
 35. Regularized Two Granularity Loss Function for Weakly Supervised Video Moment Retrieval | Request PDF - ResearchGate, accessed August 15, 2025, https://www.researchgate.net/publication/355459793_Regularized_Two_Granularity_Loss_Function_for_Weakly_Supervised_Video_Moment_Retrieval
 36. Multi-Modal Relational Graph for Cross-Modal Video Moment Retrieval - CVF Open Access, accessed August 15, 2025, https://openaccess.thecvf.com/content/CVPR2021/papers/Zeng_Multi-Modal_Relational_Graph_for_Cross-Modal_Video_Moment_Retrieval_CVPR_2021_paper.pdf
 37. Vision-Text Representation Learning for Video Moment Retrieval - QMRO Home, accessed August 15, 2025, <https://qmro.qmul.ac.uk/xmlui/handle/123456789/105857>
 38. [2312.12155] Towards Balanced Alignment: Modal-Enhanced Semantic Modeling for Video Moment Retrieval - arXiv, accessed August 15, 2025, <https://arxiv.org/abs/2312.12155>
 39. Frame-Wise Cross-Modal Matching for Video Moment Retrieval - ResearchGate, accessed August 15, 2025, https://www.researchgate.net/publication/349986007_Frame-wise_Cross-modal_Matching_for_Video_Moment_Retrieval
 40. Fewer Steps, Better Performance: Efficient Cross-Modal Clip Trimming for Video Moment Retrieval Using Language - AAAI Publications, accessed August 15, 2025, <https://ojs.aaai.org/index.php/AAAI/article/view/27941/27902>
 41. Needle in a haystack: Coarse-to-fine alignment network for moment retrieval from large-scale video collections - PMC, accessed August 15, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12080789/>
 42. Background-Aware Moment Detection for Video Moment Retrieval - CVF Open Access, accessed August 15, 2025, https://openaccess.thecvf.com/content/WACV2025/papers/Jung_Background-Aware_Moment_Detection_for_Video_Moment_Retrieval_WACV_2025_paper.pdf
 43. MomentDiff: Generative Video Moment Retrieval from Random to Real, accessed August 15, 2025, <https://arxiv.org/pdf/2307.02869>
 44. Semantic Fusion Augmentation and Semantic ... - CVF Open Access, accessed August 15, 2025, https://openaccess.thecvf.com/content/WACV2024/papers/Huang_Semantic_Fusion_Augmentation_and_Semantic_Boundary_Detection_A_Novel_Approach_WACV_2024_paper.pdf
 45. Video Moment Retrieval to Capture the Highlighting Moments in Video (Part.3) - Medium, accessed August 15, 2025, <https://medium.com/@myschang/video-moment-retrieval-to-capture-the-highlighting-moments-in-video-part-3-5222160c6fad>
 46. Zero-shot Video Moment Retrieval With Off-the-Shelf Models - Proceedings of

- Machine Learning Research, accessed August 15, 2025,
<https://proceedings.mlr.press/v203/diwan23a/diwan23a.pdf>
47. Video Moment Retrieval AI: System Architecture and Contrastive Learning - SERP AI, accessed August 15, 2025, <https://serp.ai/posts/rpus-video-moment-retrieval/>
 48. (PDF) Temporal Map-Based Boundary Refinement Network for ..., accessed August 15, 2025,
https://www.researchgate.net/publication/390964467_Temporal_Map-Based_Boundary_Refinement_Network_for_Video_Moment_Localization
 49. Moment of Untruth: Dealing with Negative Queries in Video Moment Retrieval - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2502.08544v2>
 50. TMR: Text-to-Motion Retrieval Using Contrastive 3D Human Motion Synthesis - CVF Open Access, accessed August 15, 2025,
https://openaccess.thecvf.com/content/ICCV2023/papers/Petrovich_TMR_Text-to-Motion_Retrieval_Using_Contrastive_3D_Human_Motion_Synthesis_ICCV_2023_paper.pdf
 51. Learning to Diversify for Robust Video Moment Retrieval - ResearchGate, accessed August 15, 2025,
https://www.researchgate.net/publication/385836360_Learning_to_Diversify_for_Robust_Video_Moment_Retrieval
 52. Weakly-Supervised Video Moment Retrieval via Semantic Completion Network - AAAI, accessed August 15, 2025,
<https://cdn.aaai.org/ojs/6820/6820-13-10049-1-10-20200524.pdf>
 53. Cross-Modal Video Moment Retrieval with Spatial and Language-Temporal Attention | Request PDF - ResearchGate, accessed August 15, 2025,
https://www.researchgate.net/publication/333699231_Cross-Modal_Video_Moment_Retrieval_with_Spatial_and_Language-Temporal_Attention
 54. jayleicn/moment_detr: [NeurIPS 2021] Moment-DETR code ... - GitHub, accessed August 15, 2025, https://github.com/jayleicn/moment_detr
 55. TR-DETR: Task-Reciprocal Transformer for Joint Moment Retrieval and Highlight Detection, accessed August 15, 2025, <https://arxiv.org/html/2401.02309v2>
 56. LD-DETR: Loop Decoder DETection TRansformer for Video Moment Retrieval and Highlight Detection - arXiv, accessed August 15, 2025,
<https://arxiv.org/html/2501.10787v1>
 57. MS-DETR: Towards Effective Video Moment Retrieval and Highlight Detection by Joint Motion-Semantic Learning - arXiv, accessed August 15, 2025,
<https://arxiv.org/html/2507.12062v1>
 58. MomentMix Augmentation with Length-Aware DETR for Temporally Robust Moment Retrieval - arXiv, accessed August 15, 2025,
<https://arxiv.org/html/2412.20816v2>
 59. Length-Aware DETR for Robust Moment Retrieval - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2412.20816v1>
 60. Self-Feedback DETR for Temporal Action Detection - CVF Open Access, accessed August 15, 2025,
https://openaccess.thecvf.com/content/ICCV2023/papers/Kim_Self-Feedback_DETR_for_Temporal_Action_Detection_ICCV_2023_paper.pdf

61. Dual DETRs for Multi-Label Temporal Action Detection - CVF Open Access, accessed August 15, 2025, https://openaccess.thecvf.com/content/CVPR2024/papers/Zhu_Dual_DETRs_for_Multi-Label_Temporal_Action_Detection_CVPR_2024_paper.pdf
62. [2404.00653] Dual DETRs for Multi-Label Temporal Action Detection - arXiv, accessed August 15, 2025, <https://arxiv.org/abs/2404.00653>
63. CVPR Poster Dual DETRs for Multi-Label Temporal Action Detection, accessed August 15, 2025, <https://cvpr.thecvf.com/virtual/2024/poster/30591>
64. Dual DETRs for Multi-Label Temporal Action Detection - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2404.00653v1>
65. arXiv:2404.00653v1 [cs.CV] 31 Mar 2024, accessed August 15, 2025, <https://arxiv.org/pdf/2404.00653>
66. stackoverflow.com, accessed August 15, 2025, <https://stackoverflow.com/questions/20818230/whats-the-difference-between-spatial-and-temporal-characterization-in-terms-of#:~:text=Temporal%20characterization%20occurs%20when%20you,you%20are%20analyzing%20one%20image>
67. what's the difference between spatial and temporal characterization in terms of image processing? - Stack Overflow, accessed August 15, 2025, <https://stackoverflow.com/questions/20818230/whats-the-difference-between-spatial-and-temporal-characterization-in-terms-of>
68. Learning to combine the modalities of language and video for temporal moment localization, accessed August 15, 2025, <https://ksp.etri.re.kr/ksp/article/file/65500.pdf>
69. (PDF) Where and When to Look? Spatio-temporal Attention for Action Recognition in Videos, accessed August 15, 2025, https://www.researchgate.net/publication/328212894_Where_and_When_to_Look_Spatio-temporal_Attention_for_Action_Recognition_in_Videos
70. Minimal videos: Trade-off between spatial and temporal information in human and machine vision - PMC, accessed August 15, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC7330814/>
71. [2405.12540] Context-Enhanced Video Moment Retrieval with Large Language Models, accessed August 15, 2025, <https://arxiv.org/abs/2405.12540>
72. Context-Enhanced Video Moment Retrieval with Large Language ..., accessed August 15, 2025, <https://research-repository.uwa.edu.au/files/368425736/2405.12540v1.pdf>
73. Multimodal Latent Representation Learning for Video Moment Retrieval - PMC, accessed August 15, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12299152/>
74. InstaGen: Enhancing Object Detection by Training on Synthetic Dataset - CVPR 2024 Open Access Repository, accessed August 15, 2025, https://openaccess.thecvf.com/content/CVPR2024/html/Feng_InstaGen_Enhancing_Object_Detection_by_Training_on_Synthetic_Dataset_CVPR_2024_paper.html
75. GENERALIZED VIDEO MOMENT RETRIEVAL - OpenReview, accessed August 15, 2025, <https://openreview.net/pdf?id=qdOlkeZ5e4>
76. (PDF) Moment Detection in Long Tutorial Videos - ResearchGate, accessed

August 15, 2025,

https://www.researchgate.net/publication/377432695_Moment_Detection_in_Long_Tutorial_Videos

77. Episodic Memory Representation for Long-form Video Understanding - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2508.09486v1>
78. ECCV Poster EA-VTR: Event-Aware Video-Text Retrieval, accessed August 15, 2025, <https://eccv.ecva.net/virtual/2024/poster/730>
79. Towards Efficient and Robust Moment Retrieval System: A Unified Framework for Multi-Granularity Models and Temporal Reranking - arXiv, accessed August 15, 2025, <https://arxiv.org/html/2504.08384v1>
80. Not All Inputs Are Valid: Towards Open-Set Video Moment Retrieval using Language, accessed August 15, 2025, [https://openreview.net/forum?id=h6mlfgzB44&referrer=%5Bthe%20profile%20of%20Pan%20Zhou%5D\(%2Fprofile%3Fid%3D~Pan_Zhou5\)](https://openreview.net/forum?id=h6mlfgzB44&referrer=%5Bthe%20profile%20of%20Pan%20Zhou%5D(%2Fprofile%3Fid%3D~Pan_Zhou5))
81. Improving Video Moment Retrieval by Auxiliary Moment-Query Pairs With Hyper-Interaction, accessed August 15, 2025, https://www.researchgate.net/publication/386591850_Improving_Video_Moment_Retrieval_by_Auxiliary_Moment-Query_Pairs_with_Hyper-Interaction
82. MVMR: A New Framework for Evaluating Faithfulness of Video Moment Retrieval against Multiple Distractors - Adobe Research, accessed August 15, 2025, <https://research.adobe.com/publication/mvmr-a-new-framework-for-evaluating-faithfulness-of-video-moment-retrieval-against-multiple-distractors/>